

Hands-on training for PIConGPU users

getting started with PIConGPU on LUMI

Richard Pausch¹, Sergey Bastrakov¹, Michael Bussmann^{1,2}, Finn-Ole Carstens¹, Alexander Debus¹, Fabia Dietrich¹, Sergey Ermakov^{1,3}, Julian Lenz^{1,2}, Brian Marre¹, Tapish Narwal^{1,2}, Pavel Ordyna¹, Klaus Steiniger^{1,2}, Jessica Tiebel¹, René Widera¹

¹*Helmholtz-Zentrum Dresden – Rossendorf*

²*CASUS - Center for Advanced Systems Understanding*

³*ETH Zürich*

2024-10-25

What will we do today?

- 1 check your access to LUMI
- 2 setup and run first PIConGPU simulation
- 3 data analysis of simulation using openPMD and Jupyter
- 4 dive deeper into how to setup a PIConGPU simulation
- 5 time for questions
- 6 start another project on your own



Part I

Check LUMI access

Accessing LUMI

Checking you are part of the workshop project

- go to: <https://my.lumi-supercomputer.eu/>
- login with your credentials
- go to projects and check whether you have access to **ENCCS training (EHPC-DEV-2024D09-003)**

The screenshot shows the LUMI web interface. On the left sidebar, the 'Projects' tab is selected and highlighted with a red circle. The main content area displays a table of projects. The table has columns for Name, Organization, Resources, End date, Created, and Cost estimation. The project 'ENCCS training (EHPC-DEV-2024D09-003)' is highlighted with a red underline.

Name	Organization	Resources	End date	Created	Cost estimation
Benchmarking and improving exascale in-memory streaming workflows (EHPC-BEN-2023B03-023)	Helmholtz-Zentrum Dresden-Rossendorf	0	2023-08-31	2023-05-10	0.00
<u>ENCCS training (EHPC-DEV-2024D09-003)</u>	RISE Research Institutes of Sweden AB	1	2025-09-30	2024-09-09	0.00
Plasma-PEPSC: Advancing Extreme-Scale Plasma Simulations Performance on EuroHPC Systems (EHPC-DEV-2023D12-066)	KTH 4	1	2025-03-01	2024-02-21	0.00

Accessing LUMI

Setting up an SSH key pair

- create a (LUMI specific) key pair:

```
ssh-keygen -t ed25519 -f /home/username/.ssh/id_juwels_ed25519
```

- now you should have:

- your private key: id_lumi_ed25519
- and your public key: id_lumi_ed25519.pub

- go to: <https://mms.myaccessid.org/profile/>

User profile

MENU

My profile

My linked accounts

Settings

Settings

[SSH keys](#)

- upload your **public** key


Accessing LUMI

ssh into LUMI

- execute:

```
ssh -i ~/.ssh/id_lumi_ed25519 your_username@lumi.csc.fi
```

- you should see:



Accessing LUMI

your workshop project

- check which projects you have access to:

lumi-allocations

Data updated: 2024-10-21 17:03:48							
Project	CPU (used/allocated)			GPU (used/allocated)			Storage (used/allocated)
project_465001131	0	1000	(0.0%) core/hours	5	40000	(0.0%) gpu/hours	0/90000 (0.0%) TB/hours
project_465001310	113	512000	(0.0%) core/hours	24	18000	(0.1%) gpu/hours	145/90000 (0.2%) TB/hours

- compare to:

The screenshot shows the LUMI web interface. On the left is a sidebar with navigation links: Organizations, Projects, Resources (with a sub-link for LUMI), QC, and Service catalog. The main area displays a table of project allocations. The table has columns for Name, Offering, Organization, Project, State, Created at, and Actions. Three projects are listed: 'rise-res-1-encss-tr-3-lumi-eur-2-1', 'EHPHC-DEV-2023D12-066', and 'EHPHC-BEN-2023B03-023'. Below the table, there are links for 'Plan details' and 'Attributes' for the first project, and 'Show details' for the others. The interface also includes a search bar, a sidebar with a '+ Add resource' button, and a top bar with user information (Hello Richard) and buttons for 'Import' and 'Add'.

Name	Offering	Organization	Project	State	Created at	Actions
rise-res-1-encss-tr-3-lumi-eur-2-1	LUMI EUROHPC-JU / Development Access	RISE Research Institutes of Sweden AB	ENCCS training (EHPC-DEV-2024D09-003)	OK	2024-09-09 07:21	
EHPHC-DEV-2023D12-066	LUMI EUROHPC-JU / Development Access	KTH 4	Plasma-PEPSC: Advancing Extreme-Scale Plasma Simulations Performance on EuroHPC Systems (EHPC-DEV-2023D12-066)	OK	2024-04-23 13:53	
EHPHC-BEN-2023B03-023	LUMI EUROHPC-JU / Benchmark Access	Helmholtz-Zentrum Dresden-Rossendorf	Benchmarking and improving exascale in-memory streaming workflows (EHPC-BEN-2023B03-023)	TERMINATING	2023-05-10 12:50	

Accessing LUMI

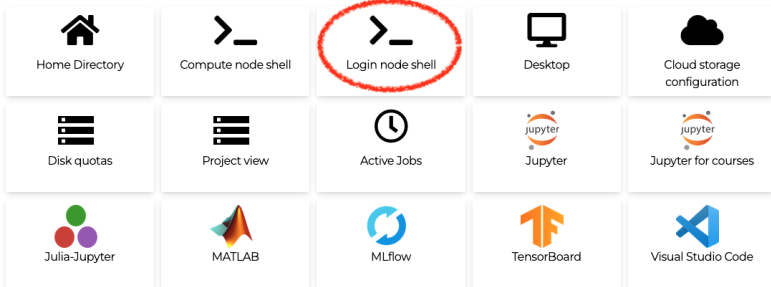
alternative access via browser

- LUMI offers a web interface

www.lumi.csc.fi

- login with your credentials
- select compute node shell

Pinned Apps



Accessing LUMI

creating your working directories

- we are working with

`project_465001310`

- project directory

`/projappl/project_465001310`

- please create a directory there

```
cd /projappl/project_465001310
mkdir $(whoami)
```

- simulation directory

`/scratch/project_465001310`

- please create a directory there

```
cd /scratch/project_465001310
mkdir $(whoami)
```



Part II

First PIConGPU simulation

Setting up PIConGPU (on LUMI)

There is a manual!

`https://picongpu.readthedocs.io`

is your first source of help!

The steps we follow now are similar to the 'PIConGPU in 5 Minutes on Hemera' tutorial there.

Performing PIconGPU simulations

The general workflow

1 Get PIconGPU

```
$ git clone https://github.com/ComputationalRadiationPhysics/picongpu.git
```

2 Prepare and load environment

```
$ source picongpu.profile
```

3 Create simulation setup

```
$ pic-create $PIC_EXAMPLES/LaserWakefield myLWFA
```

4 Adjust setup parameters

```
$ vim simulation.param
```

5 Compile setup

```
$ pic-build
```

6 Run setup (submit to batch system)

```
$ tbg -s -t -c etc/picongpu/*.cfg mySimOutputDirectory
```

7 Analyze results

```
$ jupyter-lab
```

Setting up PIConGPU on LUMI

(1) get the source code

- download the PIConGPU source code:

in home directory

```
cd $HOME/  
mkdir src  
cd src
```

- get source code

```
git clone https://github.com/ComputationalRadiationPhysics/picongpu.git
```

Setting up PICongGPU on LUMI

(2) prepare and load environment

- copy default setup file:

```
cp picongpu/etc/picongpu/lumi-eurohpc/lumi-G_hipcc_picongpu.profile.example  
$HOME/lumi-G_hipcc_picongpu.profile
```

- either use your editor of choice in the terminal (vim, emacs, nano) or use the web interface
- adjust the following lines:

```
8 export MY_MAILNOTIFY="ALL"  
9 export MY_MAIL="your.mail@server.gov"  
16 export PROJID=project_465001310  
17 export PROJECT_DIR=/projappl/$PROJID/  
18 export SCRATCH=/scratch/$PROJID/  
27 export EDITOR="emacs -nw"  
72 export PIC_LIBS=$PROJECT_DIR/workshop_software/local  
81 export PICSRC=$HOME/src/picongpu
```

```
source $HOME/lumi-G_hipcc_picongpu.profile
```

Setting up PIconGPU on LUMI

(3) Create simulation setup

```
cd $HOME  
mkdir -p picInputs  
cd picInputs  
pic-create LaserWakefield LWFA  
cd LWFA
```

■ ...

■ peek into include/picongpu/param:

```
$ ls -l include/picongpu/param
```

⇒ The *.param files largely define the simulation setup

Setting up PIConGPU on LUMI

(4.1) Adjust setup parameters

- for now we will not go into details
- just to follow the workflow, we will adjust **one compile-time** parameter
- compile-time parameters are in `include/picongpu/param`
- let's see where the parameter files are:

```
cd include/picongpu/param  
ls -l
```

- let's edit the laser configuration

```
vim incidentField.param
```

66 change a_0 from 8.0 to 5.0

Setting up PIConGPU on LUMI

(4.2) Adjust setup parameters

- just to follow the workflow, we will adjust **one run-time** parameter
- run-time parameters are in `etc/picongpu/*.cfg`
- let's see where the parameter files are:

```
cd $HOME/picInputs/LWFA
cd etc/picongpu/
ls -l
```

- let's edit the `4.cfg`

```
vim 4.cfg
```

68 remove checkpointing and add mapped memory strategy

```
TBG_openPMD="--openPMD.period 100 \\  
            --openPMD.file simData \\  
            --openPMD.ext bp \\  
            --openPMD.dataPreparationStrategy mappedMemory"
```

Setting up PIConGPU on LUMI

(5) compile setup

- go back to the root of your LWFA example

```
cd $HOME/picInputs/LWFA
```

- get compute resources for compiling:

```
getDevice
```

- compile your setup

```
pic-build
```

- better:


```
pic-build -f 2> err.txt
```

- Allows to read errors in `err.txt` (after compiling)
- Using option `-f` ensures we do not use cached `*.param` files but the original ones during compile. Accidentally using cached files is a common source of errors.
- `-c` ... fixes a current issue with CMake

Setting up PIConGPU on LUMI

You can spend compile time to rethink your parameter choices
(or check the latest xkcd...)

```
pic-build -f 2> err.txt
```



PARAMETERBALL IS A RAQUET GAME DIVIDED INTO
FOUR QUARTERS, WITH BALL SIZE, COURT SIZE,
AND NET HEIGHT RANDOMIZED EACH QUARTER.

Setting up PIConGPU on LUMI

We have a reservation

- in order to get your simulation to start, please use today's reservation
- the reservation is called:

ENCCS_day3

- adjust or tbg template file

```
vim etc/picongpu/lumi-eurohpc/standard-g.tpl
```

replace line 25:

```
#SBATCH --partition=!TBG_queue
```

with:

```
#SBATCH --reservation=ENCCS_day3
```

```
#SBATCH --exclusive
```

Setting up PIconGPU on LUMI

(6) Run setup

- **go back to LUMI head node - do not run tbg from a compute node**
- submit job using tbg (assuming you changed the standard-g.tpl)

```
tbg -s -t -c etc/picongpu/4.cfg $SCRATCH/$(whoami)/01_LWFA
```
- tbg stand for Template Batch Generator and is our abstraction for setups for various clusters
- it sets a submit command via -s (default on LUMI: sbatch)
- it uses a template via -t (default on LUMI: standard-g.tpl)
- it uses a setup configuration
- does your job run?

```
squeue -u $USER  
scontrol show jobid <Your job's id>
```



Part III

Data analysis with openPMD and Jupyter

Analysing your PIConGPU simulation

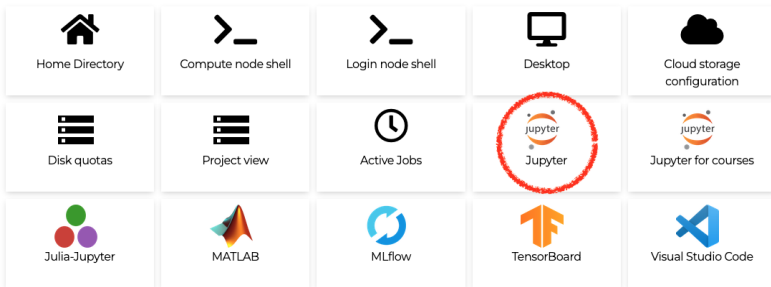
using jupyter, openPMD-api, ...

- again use the LUMI web interface

www.lumi.csc.fi

- login with your credentials
- select Jupyter

Pinned Apps



Analysing your PIConGPU simulation

getting a jupyter job

- set reservation
- use **2 cores**
- use advanced setting and use prepared python environment
- for now use your home directory as working directory

Basic

Advanced

Custom init

File

Path to script

/project/project_465001310/workshop_software/env.sh

The script on the path provided will be sourced by bash before starting Jupyter.
Ensure that the script sets up the environment so that **python** is in **PATH** and Jupyter is installed.

Launch


Analysing your PIConGPU simulation

run an example notebook

- copy notebook to your home directory

```
cp /projappl/project_465001310/workshop_software/notebooks/LWFA_analysis.ipynb  
$HOME/
```

- change the simulation directory to your simulation path



- execute it
- explore other data



Part IV

Adjusting simulation parameters

PIConGPU *.param files

`simulation.param`

- contains grid layout
 $\Delta x, \Delta y, \Delta z, \Delta t$
- reference density
BASE_DENSITY_SI
- number of particles per cell to be initialized
TYPICAL_PARTICLES_PER_CELL

PIConGPU *.param files

speciesInitialization.param

- contains a pipeline that described how macro-particles are initialized

```
using InitPipeline = pmacc::mp_list<
    ...>;
```

- particles can be created based on a density profile

```
CreateDensity<densityProfiles::Gaussian,
    startPosition::Random,
    PIC_Electrons>
```

- or derived from an already existing particle distribution

```
Derive<PIC_Electrons, PIC_Ions>
```

- particle attributes can be manipulated to set charge states, temperatures, drifts, ...

PIConGPU *.param files

density.param

- LWFA contains only a Gaussian density profile
- let's have a look at it ...
- more density profiles are available, most flexible is:

```
struct FreeFormulaFunctor
{
    HDINLINE float_X operator()(const floatD_64& position_SI,
                                const float3_64& cellSize_SI) {

        float_X s = 1.0_X - 5.0_X * math::abs(position_SI.y() - 0.0002_X);
        s *= float_X(s >= 0.0);
        return s;
    }
};

using FreeFormula = FreeFormulaImpl<FreeFormulaFunctor>;
```

PIConGPU *.param files

incidentField.param

- LWFA uses Gaussian laser profile

- at the very end:

```
using XMin = profiles::None;  
using XMax = profiles::None;  
using YMin = PARAM_LASERPROFILE;  
using YMax = profiles::None;  
using ZMin = profiles::None;  
using ZMax = profiles::None;
```

- laser definition:

```
using GaussianProfile = profiles::GaussianPulse<GaussianPulseParam>;
```

- let's have a detailed look

Task

run your own mini-simulation campaign

- based on your initial LWFA setup, either:
- simulate half and double the density and study the influence on the wakefield
- simulate $a_0 = 1.0$ and $a_0 = 0.2$ and see the effect on the wakefield
- play around with any other parameter you like



Part V

Time for your questions

- What is ...?
- How do I do ...?
- Is it possible to ...?
Yes! (with enough time and effort ;-)
- Where do I get more information on ...?

In order of precedence:

- The manual: <https://picongpu.readthedocs.io>
- Old issues
<https://github.com/ComputationalRadiationPhysics/picongpu/issues?q=is%3Aissue>
- Open a new issue



Part VI

Starting from another examples

Explore other default examples

setup your own example - we are here to help you

■ Kelvin Helmholtz instability

KelvinHelmholtz

use 4.cfg and add openPMD output

<https://github.com/ComputationalRadiationPhysics/picongpu/tree/dev/share/picongpu/examples/KelvinHelmholtz>

■ Solid density foil target for proton acceleration

FoillCT

use 4.cfg or 8.cfg

only uses 2D

<https://github.com/ComputationalRadiationPhysics/picongpu/tree/dev/share/picongpu/examples/FoillCT>