



**Plasma-PEPSC Workshop**  
**23 October 2024**

# **The Open Standard for Particle-Mesh Data**




SPONSORED BY THE



STAATSMINISTERIUM  
FÜR WISSENSCHAFT  
UND KULTUR

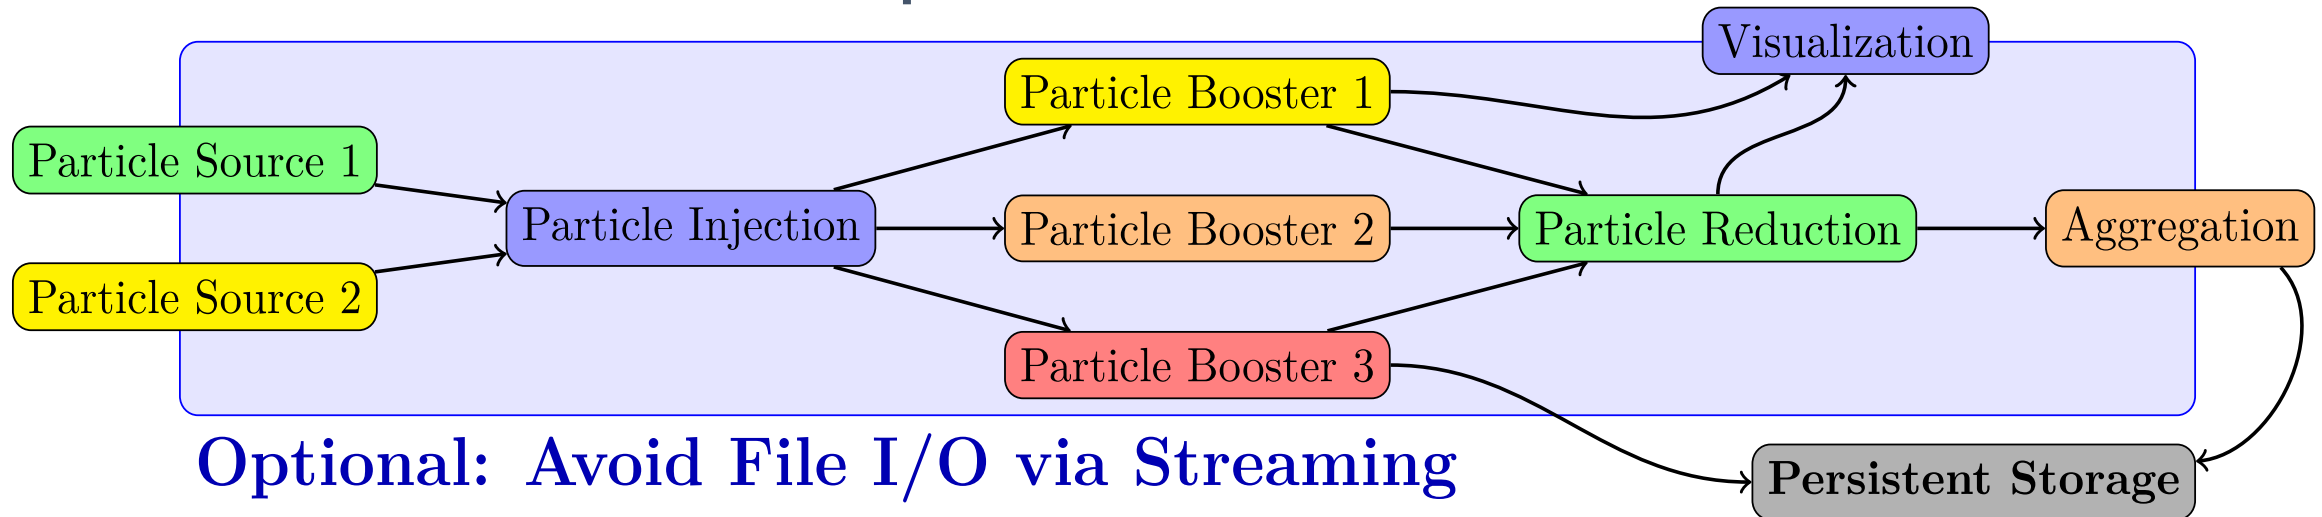


Scientific workflows are complex:




- need to span different **time** and **length scales**
- scientific modeling requires **multiple codes**, collaborating in a **data processing pipeline**
- **bridge heterogeneous models** by standardization of data


Scientific workflows are complex:




→ openPMD standard  
for **particle-mesh** data  
as communication layer



# What is particle-mesh data?



[0:3] particles   [3:6] particles   [6:10] particles



## Mesh

n-dimensional space,  
divided into discrete cells

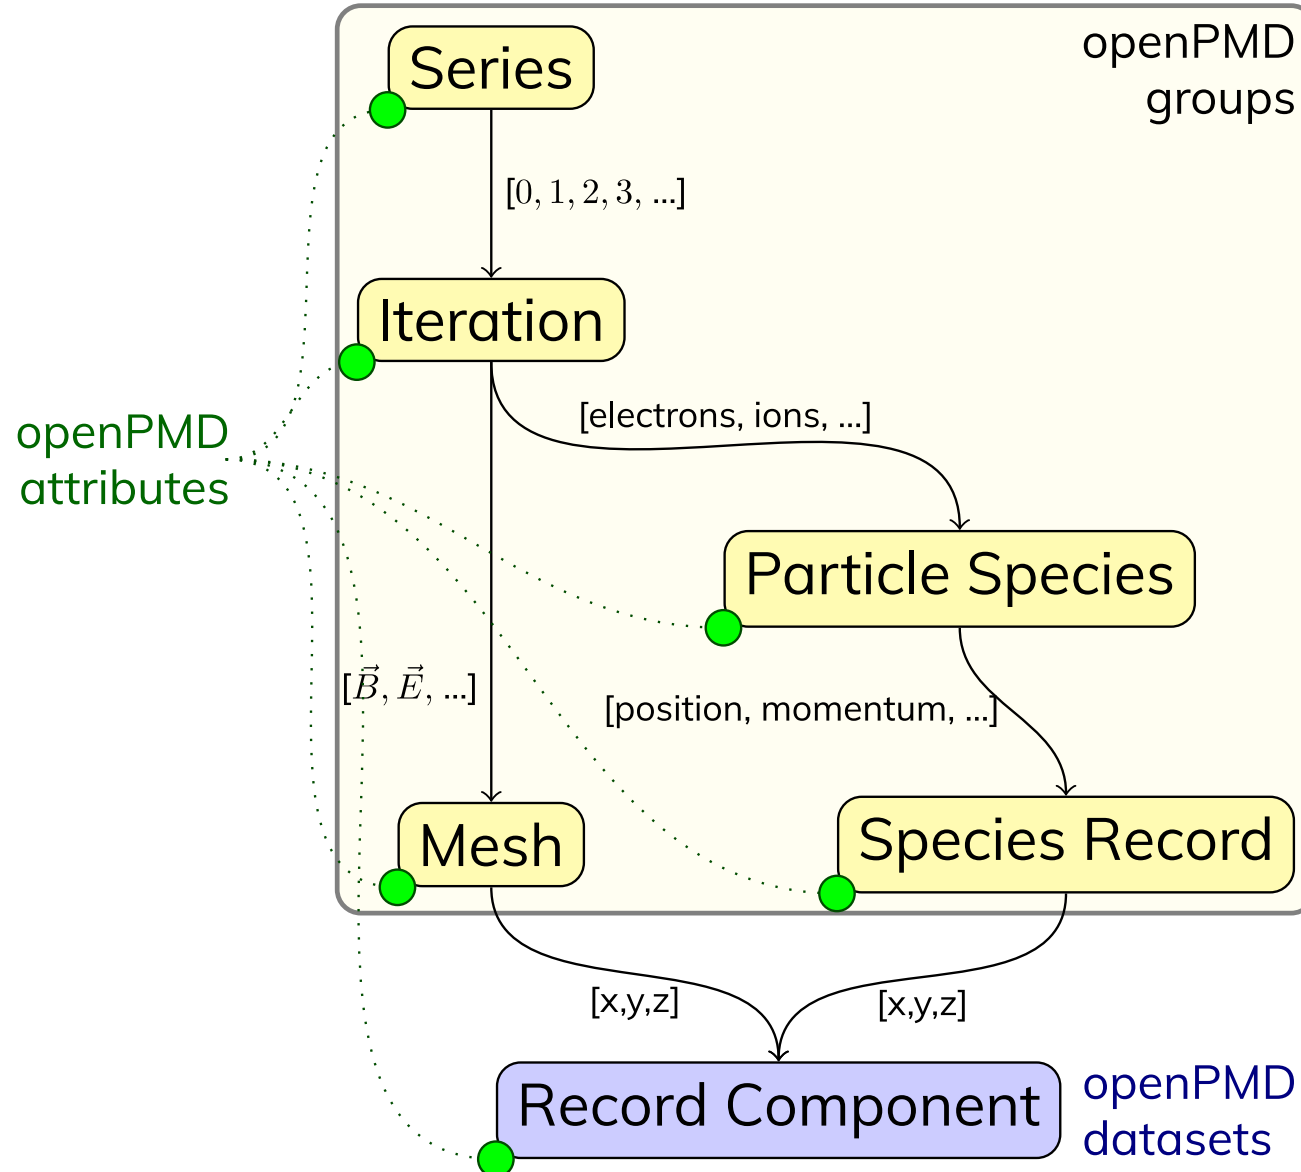
- e.g. temperature:  
store a scalar number per cell
- e.g. electrical fields:  
store a 3D vector per cell

## Particles

A list of discrete objects,  
located on the mesh

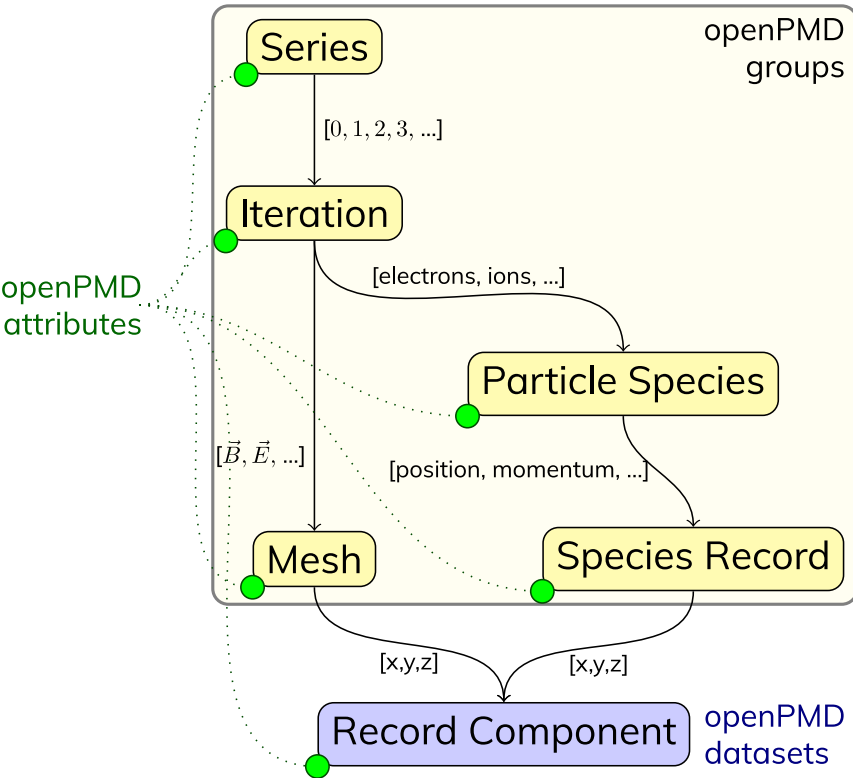
- for each particle: list its position
- optionally: list charge, weight, ...

# openPMD hierarchy



- **Structure** for series & snapshots encoded as either:
  - **files** (one file per iteration)
  - **groups** (reuse files)
  - **variables** (reuse files & variables in ADIOS2)
- Records for **physical observables** constants, mixed precision, complex numbers
- **Attributes**: unit conversion, description, relations, mesh geometry, authors, env. info, ...

# Example dataset: HDF5 backend



Sample data  
created with PIconGPU

simData\_000100.h5

- data
  - 100
    - fields
      - B
      - E
        - x
        - y
        - z
        - e\_all\_chargeDensity**
        - e\_all\_energyDensity
        - i\_all\_chargeDensity
        - i\_all\_energyDensity
    - picongpu\_idProvider
    - particles
      - e
      - i
        - charge
        - mass
        - momentum
        - particlePatches
          - extent
          - numParticles
          - numParticlesOffset
          - offset
        - position
          - x
          - y
          - z
        - positionOffset
        - weighting

Object Attribute Info    General Object Info

Attribute Creation Order:    Creation Order NOT Tracked

Number of attributes = 11

Name	Type	Array Size	Value[50](...)
axisLabels	String	3	z, y, x
dataOrder	String	Scalar	C
fieldSmoothing	String	Scalar	none
geometry	String	Scalar	cartesian
gridGlobalOffset	64-bit	3	0.0, 0.0, 0.0
gridSpacing	32-bit	3	1.7416798, 1.7416798, 1.7416798
gridUnitSI	64-bit	Scalar	5.3662849982E-8
position	32-bit	3	0.0, 0.0, 0.0
timeOffset	32-bit	Scalar	0.0
unitDimension	64-bit	7	-3.0, 0.0, 1.0, 1.0, 0.0, 0.0, 0.0
unitSI	64-bit	Scalar	338590.78364382515

# Example dataset: ADIOS2 backend

float	/data/50/fields/E/x	{128, 128, 128}
float	/data/50/fields/E/y	{128, 128, 128}
float	/data/50/fields/E/z	{128, 128, 128}
float	/data/50/particles/e/position/x	{50053105}
float	/data/50/particles/e/position/y	{50053105}
float	/data/50/particles/e/position/z	{50053105}
int32_t	/data/50/particles/e/positionOffset/x	{50053105}
int32_t	/data/50/particles/e/positionOffset/y	{50053105}
int32_t	/data/50/particles/e/positionOffset/z	{50053105}

**Hierarchical  
data organization**

***n*-dim. datasets  
for heavyweight data**

string	/data/50/fields/E/axisLabels	attr	= {"z", "y", "x"}
string	/data/50/fields/E/dataOrder	attr	= "C"
string	/data/50/fields/E/fieldSmoothing	attr	= "none"
string	/data/50/fields/E/geometry	attr	= "cartesian"
double	/data/50/fields/E/gridGlobalOffset	attr	= {0, 0, 0}
float	/data/50/fields/E/gridSpacing	attr	= {1.74168, 1.74168, 1.74168}
double	/data/50/fields/E/gridUnitSI	attr	= 5.36628e-08
float	/data/50/fields/E/timeOffset	attr	= 0
double	/data/50/fields/E/unitDimension	attr	= {1, 1, -3, -1, 0, 0, 0}
float	/data/50/fields/E/x/position	attr	= {0.5, 0, 0}
double	/data/50/fields/E/x/unitSI	attr	= 9.5224e+12
float	/data/50/fields/E/y/position	attr	= {0, 0.5, 0}
double	/data/50/fields/E/y/unitSI	attr	= 9.5224e+12
float	/data/50/fields/E/z/position	attr	= {0, 0, 0.5}
double	/data/50/fields/E/z/unitSI	attr	= 9.5224e+12

**Attributes  
for self-description**

# Example dataset: JSON/TOML backend

```
{
  "attributes": {
    "author": {
      "datatype": "STRING",
      "value": "franz"
    },
    "date": {
      "datatype": "STRING",
      "value": "2020-10-08 19:29:13 +0200"
    },
    "some more...": null
  },
  "data": {
    "0": {
      "attributes": {
        "cell_depth": {
          "datatype": "DOUBLE",
          "value": 4.252342224121094
        },
        "cell_height": {
          "datatype": "DOUBLE",
          "value": 1.0630855560302734
        },
        "cell_width": {
          "datatype": "DOUBLE",
          "value": 4.252342224121094
        },
        "many many more": null
      },
      "fields": {
        "B": {
          "attributes": {
            "axisLabels": {
              "datatype": "VEC_STRING",
              "value": [
                "z",
                "y",
                "x"
              ]
            },
            "position": {
              "datatype": "VEC_DOUBLE",
              "value": [
                0,
                0.5,
                0.5
              ]
            },
            "unitSI": {
              "datatype": "DOUBLE",
              "value": 40903.82224060171
            }
          },
          "data": [
            [
              "multidimensional dataset here"
            ]
          ]
        }
      }
    }
  }
}
```

- Part of the package: No need to install 3rd-party dependencies
- Useful for debugging and prototyping
- Limited parallel support
- Courtesy to Nils Lohmann's JSON library for C++
- With recent release: Convert output to TOML  
Idea: openPMD formatted configuration files



# Reference Implementation in C++ & Bindings: Python and Julia

Online Documentation:  
[openpmd-api.readthedocs.io](https://openpmd-api.readthedocs.io)

Open-Source Development & Tests:  
[github.com/openPMD/openPMD-api](https://github.com/openPMD/openPMD-api)

INSTALLATION

Installation

Changelog

Upgrade Guide

USAGE

Concepts

First Write

Include / Import

Open

Iteration

Attributes

Data

Record

Units

Register Chunk

Flush Chunk

After successful installation, you can start using openPMD-api as follows:

**C++17**

```
#include <openPMD/openPMD.hpp>

// example: data handling
#include <numeric> // std::iota
#include <vector>   // std::vector

namespace io = openPMD;
```

**Python**

```
import openpmd_api as io

# example: data handling
import numpy as np
```

**Open**


Write into a new openPMD series in `myOutput/data_<00...N>.h5`. Further file formats than `.h5` (HDF5) are supported: `.bp` (ADIOS1/ADIOS2) or `.json` (JSON).

```
auto series = io::Series(
    "myOutput/data_005T.h5",
    io::Access::CREATE);
```

```
series = io.Series(
    "myOutput/data_005T.h5",
    io.Access.create)
```

✓	All checks have passed	25 successful checks
✓	macOS / appleclang12_py_mpi_h5_ad2 (pull_request)	Successful in 17m <a href="#">Details</a>
✓	Windows / MSVC w/o MPI (pull_request)	Successful in 6m <a href="#">Details</a>
✓	Intel / ICC C++ only (pull_request)	Successful in 7m <a href="#">Details</a>
✓	Tooling / Clang ASAN UBSAN (pull_request)	Successful in 58m <a href="#">Details</a>
✓	Nvidia / CTK@11.2 (pull_request)	Successful in 4m <a href="#">Details</a>
✓	Linux / clang8 py38 mpich h5 ad1 ad2 newLayout (pull request)	Successful in 29m <a href="#">Details</a>

## Rapid and easy installation on any platform:



```
python3 -m pip install
openpmd-api
```



```
brew tap openpmd/openpmd
brew install openpmd-api
```



```
cmake -S . -B build
cmake --build build
--target install
```



```
conda install
-c conda-forge
openpmd-api
```



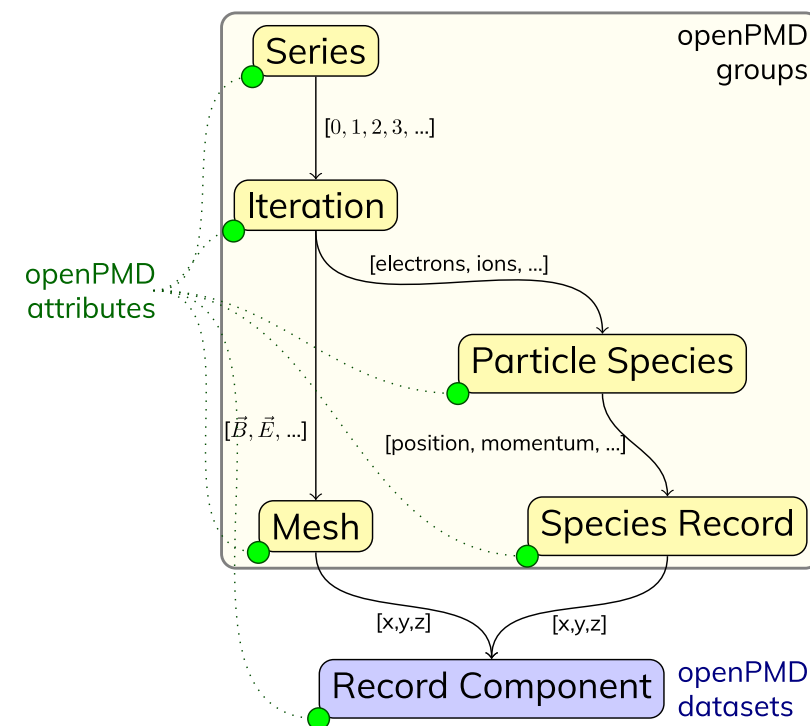
```
spack install
openpmd-api
```



```
module load openpmd-api
```

A Huebl, F Poeschel, F Koller, J Gu, et al.  
"openPMD-api: C++ & Python API for Scientific I/O with openPMD" (2018) [DOI:10.14278/rodare.27](https://doi.org/10.14278/rodare.27)

# Hands-On: openPMD-api: basic object model



**Module environment** at `/project/project_465001310/workshop_software/env.sh`

**Materials** at [https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024\\_workshop](https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop)

Read the TODO comments inside `src/openPMDOutput.hpp`:

## unitDimension


automated description of physical dimension  
only powers of base dimensions

length **L**, mass **M**, time **T**, electric current **I**,  
thermodynamic temperature **theta**,  
amount of substance **N**, luminous intensity **J**

Magnetic field:  $[B] = M / (I * T^2)$   
→ (0, 1, -2, -1, 0, 0, 0)

## unitSI (recommended)

relation to an absolute unit system




Wikimedia Commons

**Findable:** Standardized metadata to identify the data producer

```
string    /author      attr    = "franz"  
string    /software    attr    = "PIConGPU"  
string    /softwareVersion attr    = "0.5.0-dev"
```

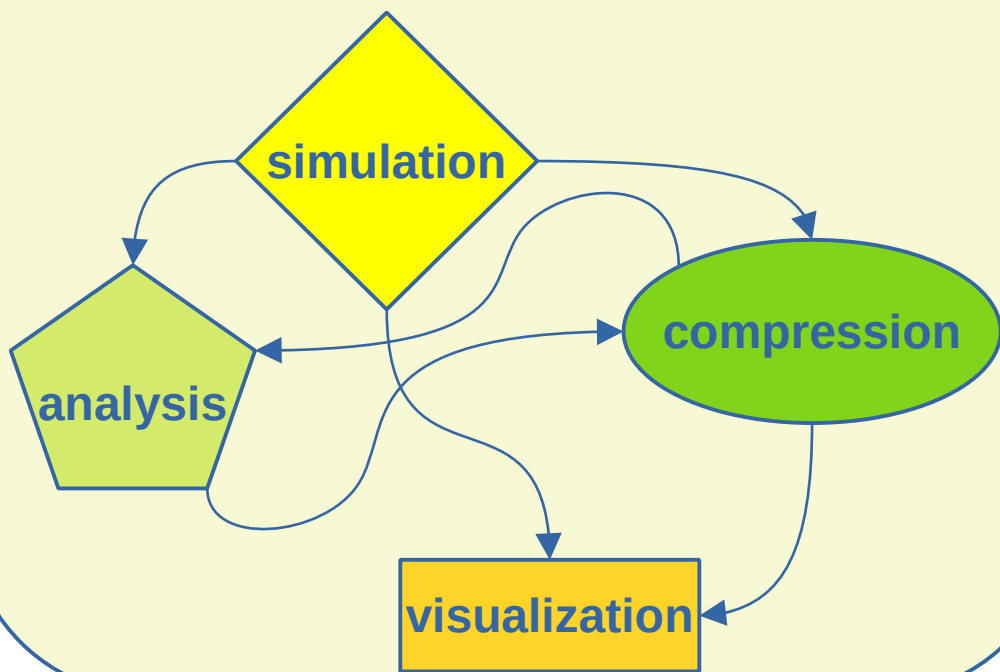
**Accessible:** Open standard, implementable in various formats



\*currently implemented,  
but not limited to

## Interoperable:

Data exchange spans applications, platforms and teams



## Reusable:


Rich and standardized description for physical quantities

Name	Value
axisLabels	[b'z' b'y' b'x']
dataOrder	b'C'
fieldSmoothing	b'none'
geometry	b'cartesian'
gridGlobalOffset	[0. 0. 0.]
gridSpacing	[4.252342 1.0630856 4.252342 ]
gridUnitSI	4.1671151662e-08
position	[0. 0. 0.]
timeOffset	0.0
unitDimension	[-3. 0. 1. 1. 0. 0. 0.]
unitSI	15399437.98944343

“The FAIR Guiding Principles for scientific data management and stewardship” (Mark D. Wilkinson et al.)

# Extensions: e.g. ED-PIC

**Field image** → field solver, smoothing



similar:

**Emittance** → particle push, field solver, shape

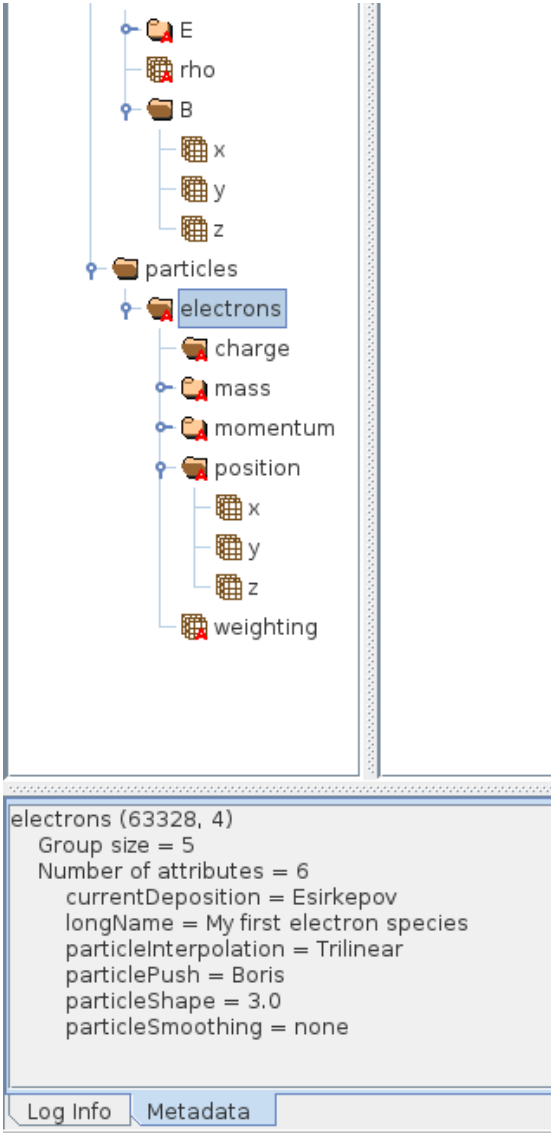


Image CC-BY 3.0: R. Lehe et al., RRSTAB 16, 021301 (2013),  
DOI:10.1103/PhysRevSTAB.16.021301

# Modeling Mesh data

```
string /data/1000/fields/E/axisLabels attr = {"z", "y", "x"}
string /data/1000/fields/E/fieldSmoothing attr = "none"
string /data/1000/fields/E/geometry attr = "cartesian"
double /data/1000/fields/E/gridGlobalOffset attr = {0, 0, 0}
float /data/1000/fields/E/gridSpacing attr = {4.25234, 1.06309, 4.25234}
double /data/1000/fields/E/gridUnitSI attr = 4.16712e-08
float /data/1000/fields/E/timeOffset attr = 0
double /data/1000/fields/E/unitDimension attr = {1, 1, -3, -1, 0, 0, 0}
float /data/1000/fields/E/x {192, 1536, 192} = 0 / 0
float /data/1000/fields/E/x/position attr = {0.5, 0, 0}
double /data/1000/fields/E/x/unitSI attr = 1.22627e+13
float /data/1000/fields/E/y {192, 1536, 192} = 0 / 0
float /data/1000/fields/E/y/position attr = {0, 0.5, 0}
double /data/1000/fields/E/y/unitSI attr = 1.22627e+13
float /data/1000/fields/E/z {192, 1536, 192} = 0 / 0
float /data/1000/fields/E/z/position attr = {0, 0, 0.5}
double /data/1000/fields/E/z/unitSI attr = 1.22627e+13
```

Mesh attributes,  
including fieldSmoothing  
defined by ED-PIC extension

x/y/z components for vector-  
type field (struct-of-array)  
including component-  
specific metadata

```
float /data/1000/fields/e_all_energyDensity {192, 1536, 192} = 0 / 0
string /data/1000/fields/e_all_energyDensity/axisLabels attr = {"z", "y", "x"}
string /data/1000/fields/e_all_energyDensity/fieldSmoothing attr = "none"
string /data/1000/fields/e_all_energyDensity/geometry attr = "cartesian"
double /data/1000/fields/e_all_energyDensity/gridGlobalOffset attr = {0, 0, 0}
float /data/1000/fields/e_all_energyDensity/gridSpacing attr = {4.25234, 1.06309, 4.25234}
double /data/1000/fields/e_all_energyDensity/gridUnitSI attr = 4.16712e-08
float /data/1000/fields/e_all_energyDensity/position attr = {0, 0, 0}
float /data/1000/fields/e_all_energyDensity/timeOffset attr = 0
double /data/1000/fields/e_all_energyDensity/unitDimension attr = {-1, 1, -2, 0, 0, 0, 0}
double /data/1000/fields/e_all_energyDensity/unitSI attr = 7.8691e+12
```

Scalar-type field  
x/y/z layer is skipped



Hands-On:

# openPMD-api: metadata


**Module environment** at `/project/project_465001310/workshop_software/env.sh`

**Materials** at [https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024\\_workshop](https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop)

Read the TODO comments inside `src/openPMDOutput.hpp`:



# openPMD-api – open stack for scientific I/O



- MPI support at all levels
- Implemented in C++17
- Bindings in C++17, Python and (dev version only) Julia
- Specify backend at runtime: I/O library, transport, compression, streaming, aggregation, ...

```
import openpmd_api as io

# pick and configure backend via JSON/TOML or inferred from filename extension
adios_config = """
    backend = "adios2"
    [[adios2.dataset.operators]]
    type = "blosc" # activate compression
"""

mode = io.Access.create
series = io.Series("simOutput.h5", mode,
    """{"hdf5": {"vfd": {"type": "subfilings"}}}""")
series = io.Series("simOutput.bp5", mode, adios_config)
series = io.Series("simOutput.sst", mode, "@./or/load/config/from/file.json")
series = io.Series("simOutput.json", mode)
```

Hands-On:

# openPMD-api: visualization, backend configuration


**Module environment** at `/project/project_465001310/workshop_software/env.sh`

**Materials** at [https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024\\_workshop](https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop)

**Launch Notebook** at <https://www.lumi.csc.fi/>

Read the instructions inside `src/next_steps.md`:

# I/O Performance lags behind Compute Performance



**Titan**

**Peak Performance:** 27 Pflop/s  
**FS Throughput:** 1 TiByte/s  
**FS Capacity:** 27 PiByte



**Summit**

200 Pflop/s  
2.5 TiByte/s  
250 PiByte



**Frontier**

1.6 Eflop/s  
5~10 TiByte/s  
500~1000 PiByte

**Growth Factor**


~60  
5~10  
18~37

- **parallel bandwidth** insufficient for HPC at full scale
- **filesystem capacity** insufficient for HPC at full scale

Same trend in **experiments**?

- Increasing **camera resolutions and data rates**

# Compute Performance Outpaces Storage Performance



**Titan**

**Peak Performance:** 27 Pflop/s  
**FS Throughput:** 1 TiByte/s  
**FS Capacity:** 27 PiByte



**Summit**

200 Pflop/s  
2.5 TiByte/s  
250 PiByte



**Frontier**

1.6 Eflop/s  
5~10 TiByte/s  
500~1000 PiByte

**Growth Factor**


~60

5~10



18~37

## Why does this concern us?

- Heterogeneous data processing pipelines traditionally have large I/O usage
- Scalable alternative: Streaming



# Streaming: Don't touch the Filesystem at all




- Data processing pipelines and increasingly experiments setups have large I/O usage
- Scalable alternative: **Streaming**  
e.g. via Infiniband (on HPC systems)  
or wide area networks (in lab settings)

## Challenge:

Compute a **balanced, aligned, local** mapping between two applications that remains useful in the problem domain

# Break through Filesystem Bandwidth with Streaming




## Memory-bound simulations reach the I/O system limits at a fraction of full scale

- Summit FS bandwidth (2.5TiByte/s) reached at 512 nodes (~11% of system size)
- Streaming workflows unaffected by filesystem bandwidth, use Infiniband hardware to scale beyond it

(benchmarks at 1024 nodes done after Summit system upgrade)

# Summit: Performance fluctuations on single ranks




## Same results, different display:

- Plot every single measurement
- Visualize reproducibility
- Box: 50% of measurement points
- Whiskers: “normal” measurements
- Others: outliers

## Evaluation for file-only setup:

- Median time slightly raised at 512 nodes  
Scaling stops due to PFS limit after that
- Outliers increasing with scale
- Outliers fatal in parallel contexts





## Same results, different display:


- Plot every single measurement
- Visualize reproducibility
- Box: 50% of measurement points
- Whiskers: “normal” measurements
- Others: outliers

## stream+file setup (stream part):

- Overall times are lower
- Median between 5 and 7 seconds
- Outliers less dominating by far



# For good throughput: Local streaming patterns, Infiniband/RDMA



## Local streaming:

Distribute data chunks only within a node

## Non-local streaming:


Distribute data chunks globally, optimize for balance and alignment

## Straight lines:

Infiniband/RDMA

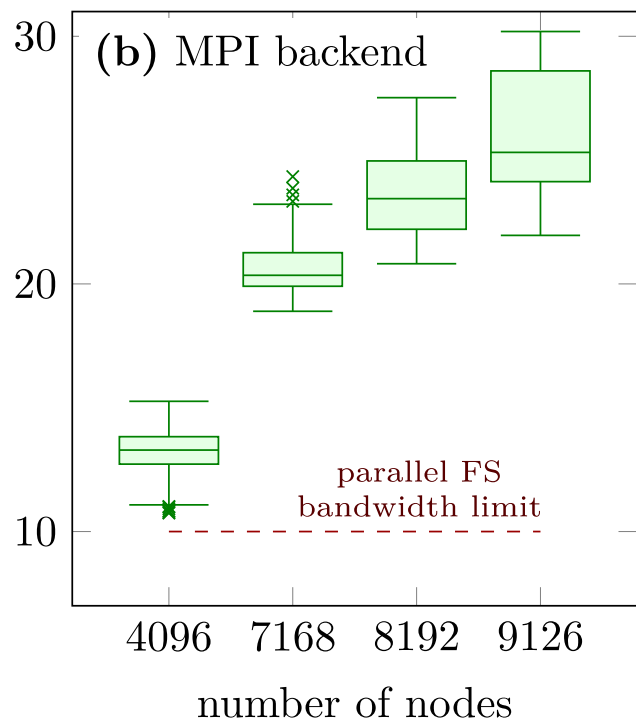
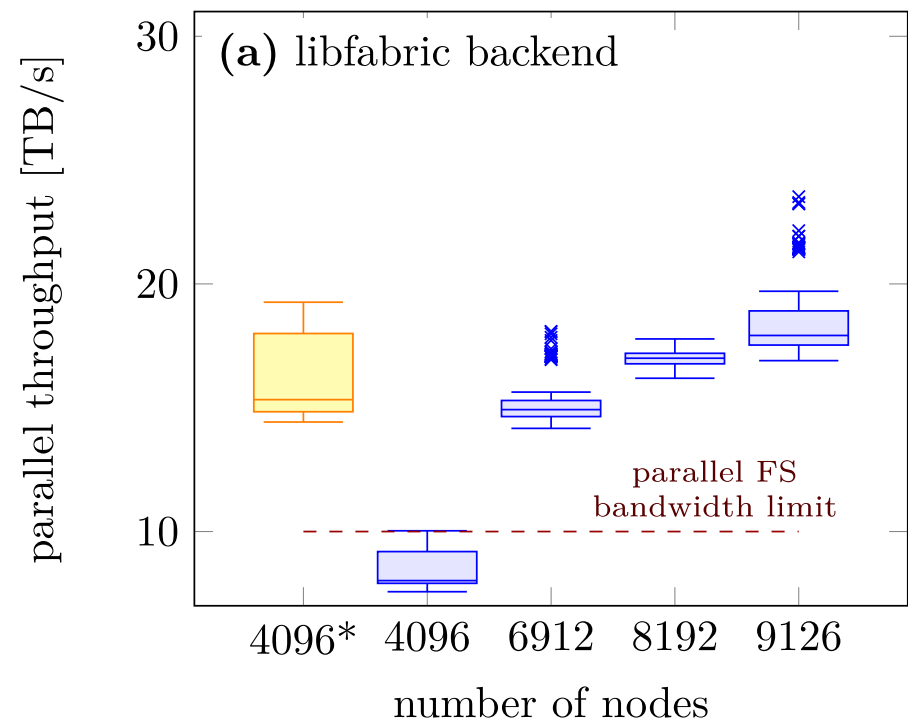
## Dashed lines:

TCP/sockets



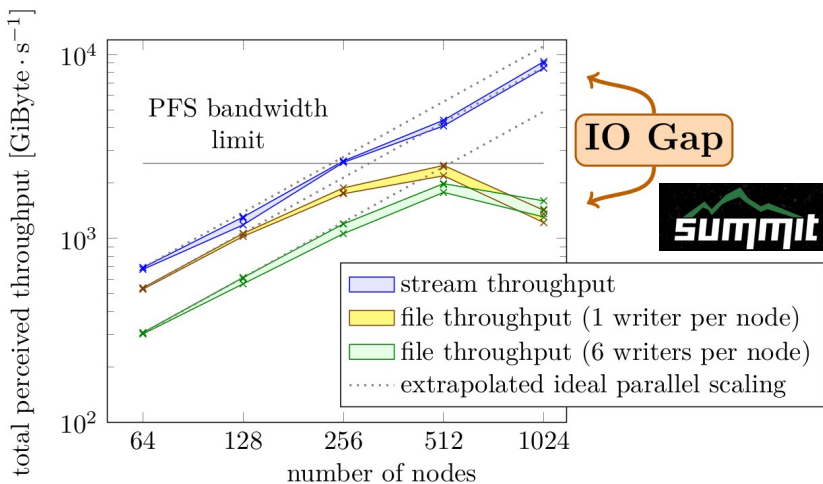
Setup: Couple PIconGPU with a scattering code (GAPD) exchange particle data only

# Break through Filesystem Bandwidth with Streaming



**Full scale of Frontier:**  
Leaving behind us  
the theoretical bandwidth  
of the parallel filesystem


# Performance: Data Layouts and no-file I/O



## Streaming Data Pipelines:

[DOI:10.1007/978-3-030-96498-6\\_6](https://doi.org/10.1007/978-3-030-96498-6_6)

by F Poeschel, A Huebl et al., SMC21 (2022)




## Fast Compressors Needed:



[DOI:10.1007/978-3-319-67630-2\\_2](https://doi.org/10.1007/978-3-319-67630-2_2)

by A Huebl et al., ISC DRBSD-1 (2017)

## ADIOS openPMD-api w/ WarpX



>5.5TB/s FS BW: two-tier lustre w/ high-performance storage & progressive files



Impact of decomposition schemes when reading

## Online Data Layout Reorganization:

[DOI:10.1109/TPDS.2021.3100784](https://doi.org/10.1109/TPDS.2021.3100784)

by L Wan, A Huebl et al., TPDS (2021)

Hands-On:

# openPMD-api: streaming I/O

**Module environment** at `/project/project_465001310/workshop_software/env.sh`

**Materials** at `https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024\_workshop`

**Launch Notebook** at `https://www.lumi.csc.fi/`

Read the instructions inside `src/visualize.py` (open as a Jupyter Notebook):

# openPMD powered Projects and Users



## Documents:

- **openPMD standard** (1.0.0, 1.0.1, 1.1.0)  
*the underlying file markup and definition*  
A Huebl et al., doi: 10.5281/zenodo.33624

## Language Binding:

- **openPMD-api** (HZDR, CASUS, LBNL)  
*reference API for openPMD data handling*  
maintainers: A Huebl, J Gu, F Poeschel et al.

**HiPACE++** → VisualPIC  
Credit: M.Thévenet  
& A. Ferran Pousa (DESY)




**MALA** → ParaView  
Credit: A. Cangi (CASUS)

- **Wake-T** (DESY)  
*fast particle-tracking code for plasma-based accelerators*  
maintainer: A Ferran Pousa
- **HiPACE++** (DESY, LBNL)  
*3D GPU-capable quasi-static PIC code for plasma accel.*  
maintainers: M Thevenet, S Diederichs, A Huebl
- **Bmad** (Cornell)  
*library for charged-particle dynamics simulations*  
maintainers: D Sagan et al.
- **MALA** (CASUS, SNL)  
*ML models that replace DFT calculations in materials science*  
maintainers: Attila Cangi & Sivasankaran Rajamanickam
- and more...

see also: <https://github.com/openPMD/openPMD-projects>

# Analysis and Visualization




openPMD/openPMD-viewer



```
In [1]: import numpy as np
import matplotlib notebook
# or '%matplotlib inline' for non-interactive plots
# or '%matplotlib widget' when using JupyterLab (github.com/matplotlib/jupyter-matplotlib)
import matplotlib.pyplot as plt
from openpmd_viewer import OpenPMDTimeSeries
```

```
In [2]: # Replace the string below, to point to your data
ts = OpenPMDTimeSeries('/home/franzpoeschel/singularity_build/pic_run/openPMD')
```

```
In [3]: # Interactive GUI
ts.slider()
```



- + iteration  1900

▼ Field type

Field:

B E e\_all\_chargeDensity

e\_all\_energyDensity picongpu\_idProvider

Coord:

x y z

► Slice selection

► Plotting options

Always refresh Refresh now!

► Particle quantities

► Particle selection

► Plotting options

Always refresh Refresh now!


**Standardization of data**

→ integration into modern scientific compute workflows



**DASK**

**ParaView**





Hands-On:

# **openPMD-viewer:** visualizing data written by a PIConGPU LaserWakefield simulation


**Module environment** at `/project/project_465001310/workshop_software/env.sh`

**Materials** at [https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024\\_workshop](https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop)

**Launch Notebook** at <https://www.lumi.csc.fi/>

Read the instructions inside `next_steps.md`:

Just compress the data and everything will be fine ...?





# Hands-On: **compression**

**Module environment** at `/project/project_465001310/workshop_software/env.sh`

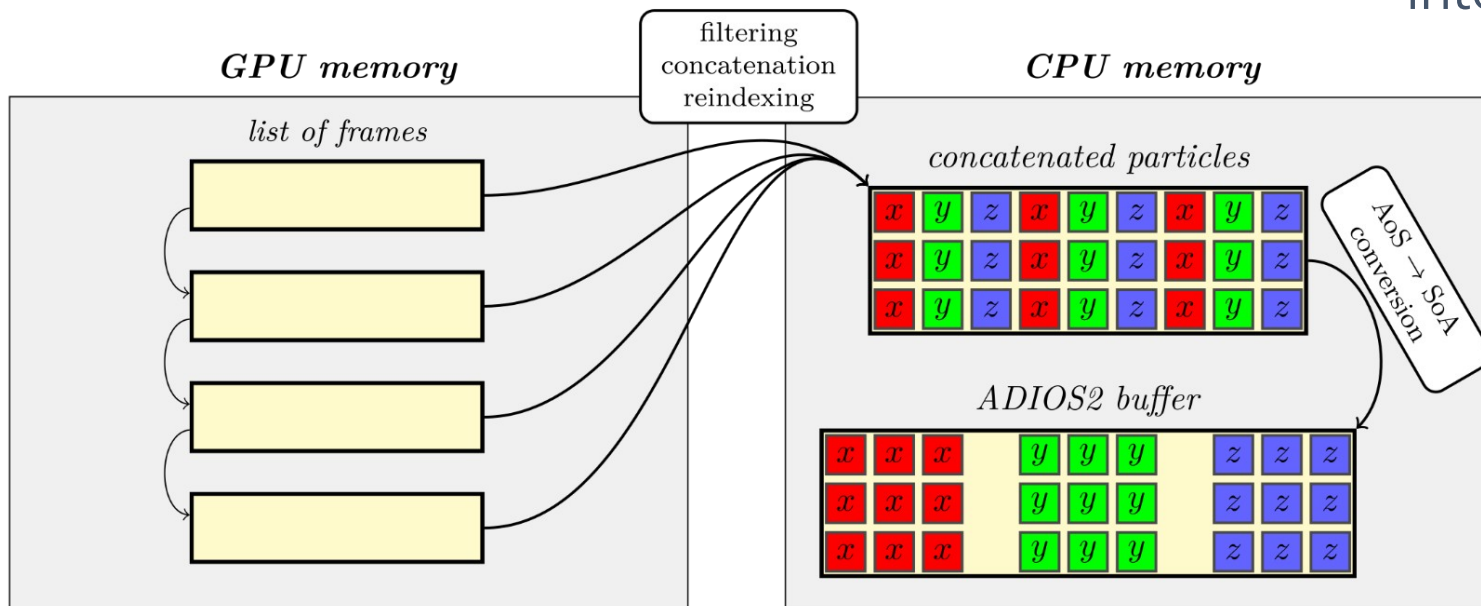
**Materials** at `https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024\_workshop`

**Launch Notebook** at `https://www.lumi.csc.fi/`

Continue with the instructions in the openPMD-viewer Notebook

# Data path in a realistic application

- In real-world applications, data is often kept in a custom, algorithm-driven format
- Reorganization needed before output
- However: ADIOS2 (optionally) buffers data before output for I/O performance
- Few big operations better than many small operations
- Elide one data copy by writing directly into those internal buffers



## Example:

Transforming  
PConGPU particle data  
for output

Advanced Hands-On:

# openPMD-api:

## Span API

**Module environment** at `/project/project_465001310/workshop_software/env.sh`

**Materials** at `https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024\_workshop`

**Launch Notebook** at `https://www.lumi.csc.fi/`

Read the instructions inside `src/openPMDOutput.hpp`

# Modeling particle data

```
uint32_t /data/particles/e/position/macroWeighted
float /data/particles/e/position/timeOffset
double /data/particles/e/position/unitDimension
double /data/particles/e/position/weightingPower
float /data/particles/e/position/x
double /data/particles/e/position/x/unitSI
float /data/particles/e/position/y
double /data/particles/e/position/y/unitSI
float /data/particles/e/position/z
double /data/particles/e/position/z/unitSI
```

```
attr = 0
attr = 0
attr = {1, 0, 0, 0, 0, 0, 0}
attr = 0
{109134176} = 0 / 0
attr = 1.772e-07
{109134176} = 0 / 0
attr = 4.43e-08
{109134176} = 0 / 0
attr = 1.772e-07
```

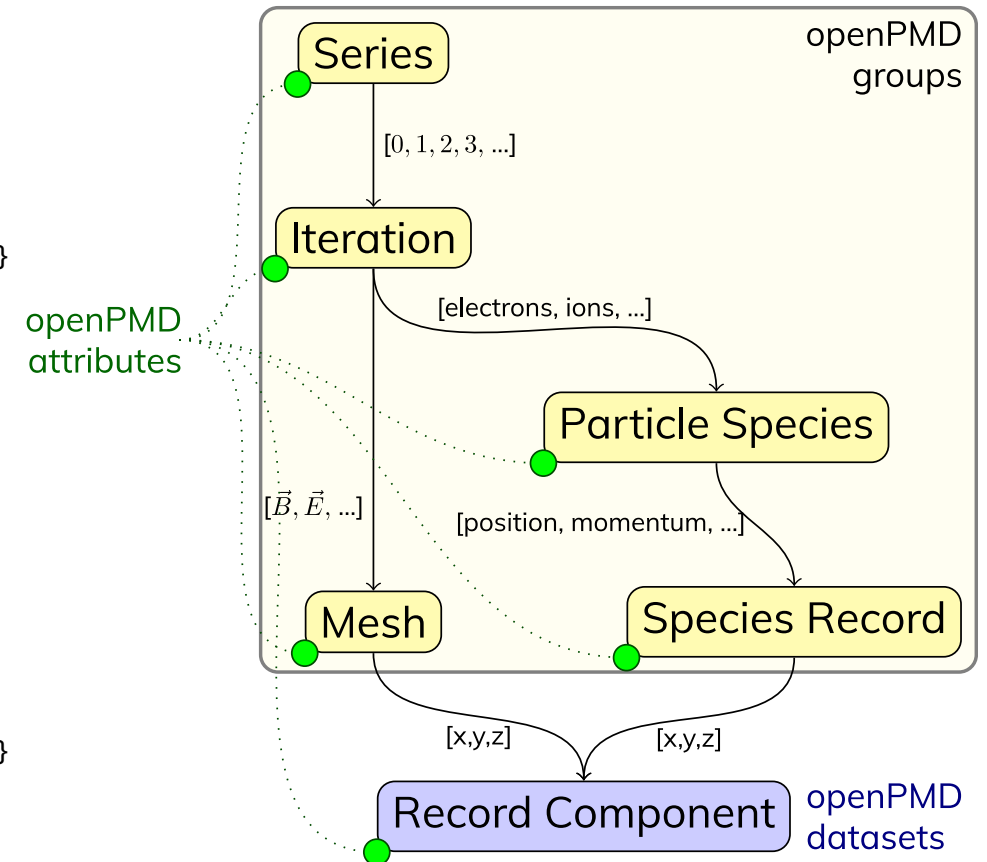
```
uint32_t /data/particles/e/positionOffset/macroWeighted
float /data/particles/e/positionOffset/timeOffset
double /data/particles/e/positionOffset/unitDimension
double /data/particles/e/positionOffset/weightingPower
int32_t /data/particles/e/positionOffset/x
double /data/particles/e/positionOffset/x/unitSI
int32_t /data/particles/e/positionOffset/y
double /data/particles/e/positionOffset/y/unitSI
int32_t /data/particles/e/positionOffset/z
double /data/particles/e/positionOffset/z/unitSI
```

```
attr = 0
attr = 0
attr = {1, 0, 0, 0, 0, 0, 0}
attr = 0
{109134176} = 0 / 0
attr = 1.772e-07
{109134176} = 0 / 0
attr = 4.43e-08
{109134176} = 0 / 0
attr = 1.772e-07
```

```
float /data/particles/e/weighting
uint32_t /data/particles/e/weighting/macroWeighted
float /data/particles/e/weighting/timeOffset
double /data/particles/e/weighting/unitDimension
double /data/particles/e/weighting/unitSI
double /data/particles/e/weighting/weightingPower
```

```
{109134176} = 0 / 0
attr = 1
attr = 0
attr = {0, 0, 0, 0, 0, 0, 0}
attr = 1
attr = 1
```

Let's have a detailed look  
at the electron species  
written by PIconGPU...



# Modeling particle data: mandatory records, vector vs. scalar

uint32_t	/data/particles/e/position/macroWeighted	attr = 0
float	/data/particles/e/position/timeOffset	attr = 0
double	/data/particles/e/position/unitDimension	attr = {1, 0, 0, 0, 0, 0, 0}
double	/data/particles/e/position/weightingPower	attr = 0
<b>float</b>	<b>/data/particles/e/position/x</b>	<b>{109134176} = 0 / 0</b>
double	/data/particles/e/position/x/unitSI	attr = 1.772e-07
<b>float</b>	<b>/data/particles/e/position/y</b>	<b>{109134176} = 0 / 0</b>
double	/data/particles/e/position/y/unitSI	attr = 4.43e-08
<b>float</b>	<b>/data/particles/e/position/z</b>	<b>{109134176} = 0 / 0</b>
double	/data/particles/e/position/z/unitSI	attr = 1.772e-07

**position:** mandatory record for particles

position is a **vector quantity**, i.e. has x/y/z  
 struct-of-array-like layout:  
 3 vectors of same length

uint32_t	/data/particles/e/positionOffset/macroWeighted	attr = 0
float	/data/particles/e/positionOffset/timeOffset	attr = 0
double	/data/particles/e/positionOffset/unitDimension	attr = {1, 0, 0, 0, 0, 0, 0}
double	/data/particles/e/positionOffset/weightingPower	attr = 0
<b>int32_t</b>	<b>/data/particles/e/positionOffset/x</b>	<b>{109134176} = 0 / 0</b>
double	/data/particles/e/positionOffset/x/unitSI	attr = 1.772e-07
<b>int32_t</b>	<b>/data/particles/e/positionOffset/y</b>	<b>{109134176} = 0 / 0</b>
double	/data/particles/e/positionOffset/y/unitSI	attr = 4.43e-08
<b>int32_t</b>	<b>/data/particles/e/positionOffset/z</b>	<b>{109134176} = 0 / 0</b>
double	/data/particles/e/positionOffset/z/unitSI	attr = 1.772e-07

**positionOffset:** also mandatory  
 specifies the base for the position  
 used for numerical reasons

<b>float</b>	<b>/data/particles/e/weighting</b>	<b>{109134176} = 0 / 0</b>
uint32_t	/data/particles/e/weighting/macroWeighted	attr = 1
float	/data/particles/e/weighting/timeOffset	attr = 0
double	/data/particles/e/weighting/unitDimension	attr = {0, 0, 0, 0, 0, 0, 0}
double	/data/particles/e/weighting/unitSI	attr = 1
double	/data/particles/e/weighting/weightingPower	attr = 1

**weighting:** custom record in PIconGPU  
 to model macro particles  
 scalar quantity, x/y/z is skipped

# Modeling particle data: constant components

string	/data/particles/e/currentDeposition	attr	= "Esirkepov"
string	/data/particles/e/particleInterpolation	attr	= "uniform"
string	/data/particles/e/particlePush	attr	= "Boris"
double	/data/particles/e/particleShape	attr	= 2
string	/data/particles/e/particleSmoothing	attr	= "none"
int32_t	/data/particles/e/charge/macroWeighted	attr	= 0
<b>uint64_t</b>	<b>/data/particles/e/charge/shape</b>	<b>attr</b>	<b>= {109134176}</b>
double	/data/particles/e/charge/timeOffset	attr	= 0
double	/data/particles/e/charge/unitDimension	attr	= {0, 0, 1, 1, 0, 0, 0}
double	/data/particles/e/charge/unitSI	attr	= 1.11432e-15
<b>double</b>	<b>/data/particles/e/charge/value</b>	<b>attr</b>	<b>= -0.00014378</b>
double	/data/particles/e/charge/weightingPower	attr	= 1
int32_t	/data/particles/e/mass/macroWeighted	attr	= 0
<b>uint64_t</b>	<b>/data/particles/e/mass/shape</b>	<b>attr</b>	<b>= {109134176}</b>
double	/data/particles/e/mass/timeOffset	attr	= 0
double	/data/particles/e/mass/unitDimension	attr	= {0, 1, 0, 0, 0, 0, 0}
double	/data/particles/e/mass/unitSI	attr	= 6.33564e-27
<b>double</b>	<b>/data/particles/e/mass/value</b>	<b>attr</b>	<b>= 0.00014378</b>
double	/data/particles/e/mass/weightingPower	attr	= 1
uint32_t	/data/particles/e/momentum/macroWeighted	attr	= 1
float	/data/particles/e/momentum/timeOffset	attr	= 0
double	/data/particles/e/momentum/unitDimension	attr	= {1, 1, -1, 0, 0, 0, 0}
double	/data/particles/e/momentum/weightingPower	attr	= 1
<b>float</b>	<b>/data/particles/e/momentum/x</b>	<b>{109134176} = 0 / 0</b>	
double	/data/particles/e/momentum/x/unitSI	attr	= 1.89938e-18
<b>float</b>	<b>/data/particles/e/momentum/y</b>	<b>{109134176} = 0 / 0</b>	
double	/data/particles/e/momentum/y/unitSI	attr	= 1.89938e-18
<b>float</b>	<b>/data/particles/e/momentum/z</b>	<b>{109134176} = 0 / 0</b>	
double	/data/particles/e/momentum/z/unitSI	attr	= 1.89938e-18

particle attributes defined  
by the ED-PIC extension

## charge and mass:

are fixed for one particle species,  
i.e. record has one constant value  
→ define via attributes shape and value  
instead of writing the whole array  
scalar quantities as well

## momentum:

again a non-constant  
vector quantity

# Modeling particle data: particle patches

```
uint64_t /data/1000/particles/e/particlePatches/offset/x {8}
(0)      0 0 0 0 96 96 96 96

uint64_t /data/1000/particles/e/particlePatches/offset/y {8}
(0)      0 512 1024 1536 0 512 1024 1536

uint64_t /data/1000/particles/e/particlePatches/offset/z {8}
(0)      0 0 0 0 0 0 0 0

uint64_t /data/1000/particles/e/particlePatches/extent/x {8}
(0)      96 96 96 96 96 96 96 96

uint64_t /data/1000/particles/e/particlePatches/extent/y {8}
(0)      512 512 512 0 512 512 512 0


uint64_t /data/1000/particles/e/particlePatches/extent/z {8}
(0)      192 192 192 192 192 192 192 192

uint64_t /data/1000/particles/e/particlePatches/numParticlesOffset {8}
(0)      0 16916087 35692896 54567264 54567264 71482787 90259808 109134176

uint64_t /data/1000/particles/e/particlePatches/numParticles {8}
(0)      16916087 18776809 18874368 0 16915523 18777021 18874368 0
```

- Common case:** Each GPU computes one sub-areal of the global simulation space  
→ Particles are written in patches, each GPU's patch is spatially bounded
- offset/extent define the box wherein the particles are located
  - numParticlesOffset/numParticles define the patch's position in the list

e.g.: GPU 0 computes the box  
(0, 0, 0)–(96, 512, 192)  
(in simulation units),  
its electron data is found  
from index 0 to 16916087  
**Note:** extent is always  
relative to the offset



Advanced Hands-On:

# openPMD-api:

## Python API: constant components

**Module environment** at `/project/project_465001310/workshop_software/env.sh`

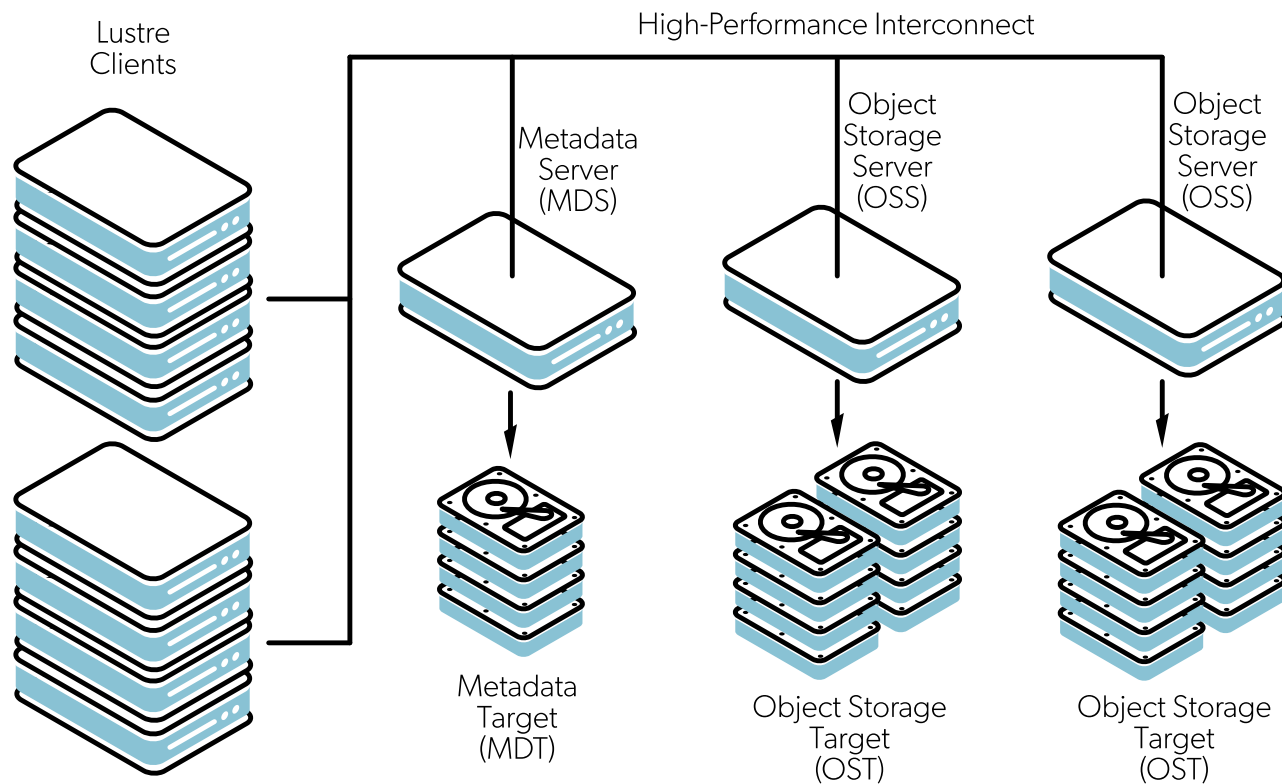
**Materials** at [https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024\\_workshop](https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop)

**Launch Notebook** at <https://www.lumi.csc.fi/>

Read the TODO comments inside `write_parallel.py`



# Using parallel filesystems

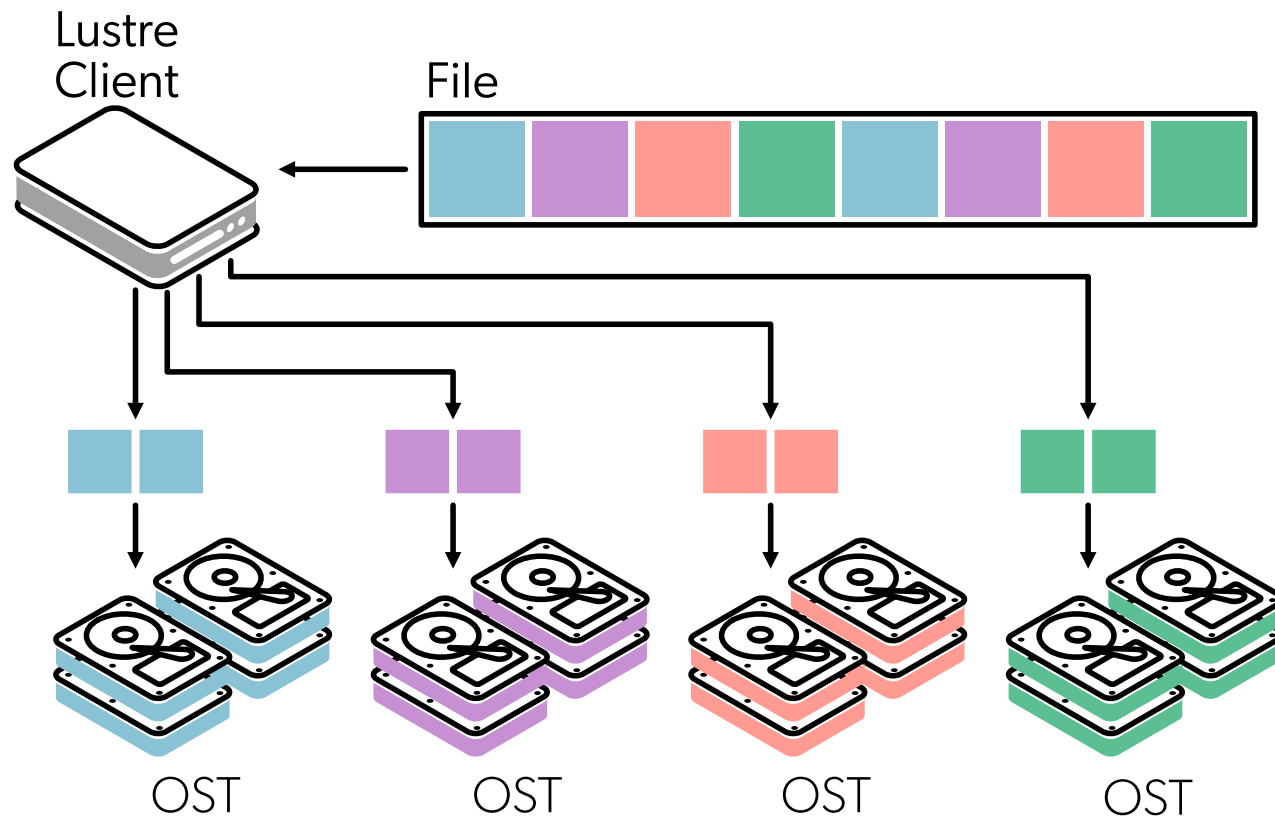


Lumi uses Lustre,  
a popular parallel filesystem

- Parallel filesystems are built from distributed hardware
- One or more metadata servers manage one or more metadata targets
- Many object storage servers manage each one or more object storage targets
- **Key to I/O performance:**  
Use the parallel resources  
(but mind that they are shared with other users)

Image from  
<https://docs.lumi-supercomputer.eu/storage/parallel-filesystems/lustre/>

# File striping – an approach at parallel I/O performance



Write file in parallel by striping it across multiple OSTs

- **stripe\_size** sets the homogeneous length of each stripe
- **stripe\_count** sets the number of used OSTs


Main use:

- Improving speed of **serial I/O**
- Improving speed of parallel I/O to a **single file**
- At low scale: Improving speed of parallel I/O to multiple files


Better: Write **multiple files** from the start already → **Subfiling**

# Subfiling – key for performance at full scale

## ADIOS openPMD-api w/ WarpX



>5.5TB/s FS BW: two-tier lustre w/ high-performance storage & progressive files



L. Wan et al.: Data Management Challenges of Exascale Scientific Simulations:  
A Case Study with the Gyrokinetic Toroidal Code and ADIOS

Subfiling is NOT independent I/O

- Mapping processes to files **N:1** scales badly
- Mapping processes to files **N:N** scales up to a mediocre size, then overloads the filesystem
- Subfiling: **N:M** mapping where  $M \leq N$
- Many different aggregation schemes in ADIOS2 and HDF5, one size does not fit all

**ADIOS2:** Aggregate to one writer per node by default

**HDF5:** Complex  $N \rightarrow N$  aggregation scheme due to structured output to disk

Advanced Hands-On:

# **openPMD-api:**

## Python API: Parallel Writing

**Module environment** at `/project/project_465001310/workshop_software/env.sh`


**Materials** at `https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024\_workshop`

**Launch Notebook** at `https://www.lumi.csc.fi/`

Read the TODO comments inside `write_parallel.py`

# Where to find us

→ head to <https://github.com/openPMD/>




## openPMD

Open Standard for Particle-Mesh Data Files

<https://www.openPMD.org> [axelhuebl@lbl.gov](mailto:axelhuebl@lbl.gov)

**Repositories** 17 Packages People 50 Teams 5 Projects


Pinned repositories



### openPMD-standard

Open Standard for Particle-Mesh Data Files


☆ 41 🍴 17



### openPMD-projects

Overview on Projects around openPMD


☆ 4 🍴 4



### openPMD-viewer

Python visualization tools for openPMD files


Jupyter Notebook ☆ 35 🍴 26



### openPMD-api

C++ & Python API for Scientific I/O


C++ ☆ 55 🍴 30



### openPMD-visit-plugin

Plugin allowing VisIt to read openPMD files

C ☆ 8 🍴 3




### openPMD-example-datasets

HDF5 Example Files

Python ☆ 5 🍴 1

...and of course <https://openpmd-api.readthedocs.io/>



The **openPMD standard** is co-authored by [Axel Huebl](#), [Rémi Lehe](#), Jean-Luc Vay, David P. Grote, Ivo F. Sbalzarini, Stephan Kuschel, David Sagan, Frédéric Pérez, Fabian Koller, [Franz Poeschel](#), Carsten Fortmann-Grote, Ángel Ferran Pousa, Juncheng E, [Maxence Thévenet](#), and Michael Bussmann.

The authors are thankful for the **community contributions** to libraries, software ecosystem, user support, review and integrations. Particularly, thank you to Yaser Afshar, Lígia Diana Amorim, James Amundson, Weiming An, Igor Andriyash, Ksenia Bastrakova, [Jean Luca Bez](#), Richard Briggs, Heiko Bureau, Jong Choi, Ray Donnelly, Dmitry Ganyushin, Marco Garten, Lixin Ge, Berk Geveci, Daniel Grassinger, Alexander Grund, [Junmin Gu](#), Marc W. Guetg, Ulrik Günther, Sören Jalas, Manuel Kirchen, John Kirkham, Scott Klasky, Noah Klemm, Fabian Koller, Mathieu Lobet, Christopher Mayes, Ritiek Malhotra, Paweł Ordyna, Richard Pausch, [Norbert Podhorszki](#), David Pugmire, Felix Schmitt, [Erik Schnetter](#), Dominik Stańczak, Klaus Steiniger, Michael Sippel, Frank Tsung, Lipeng Wan, René Widera, and Erik Zenker!