




Plasma-PEPSC Workshop
23 October 2024


The Open Standard for Particle-Mesh Data



SPONSORED BY THE

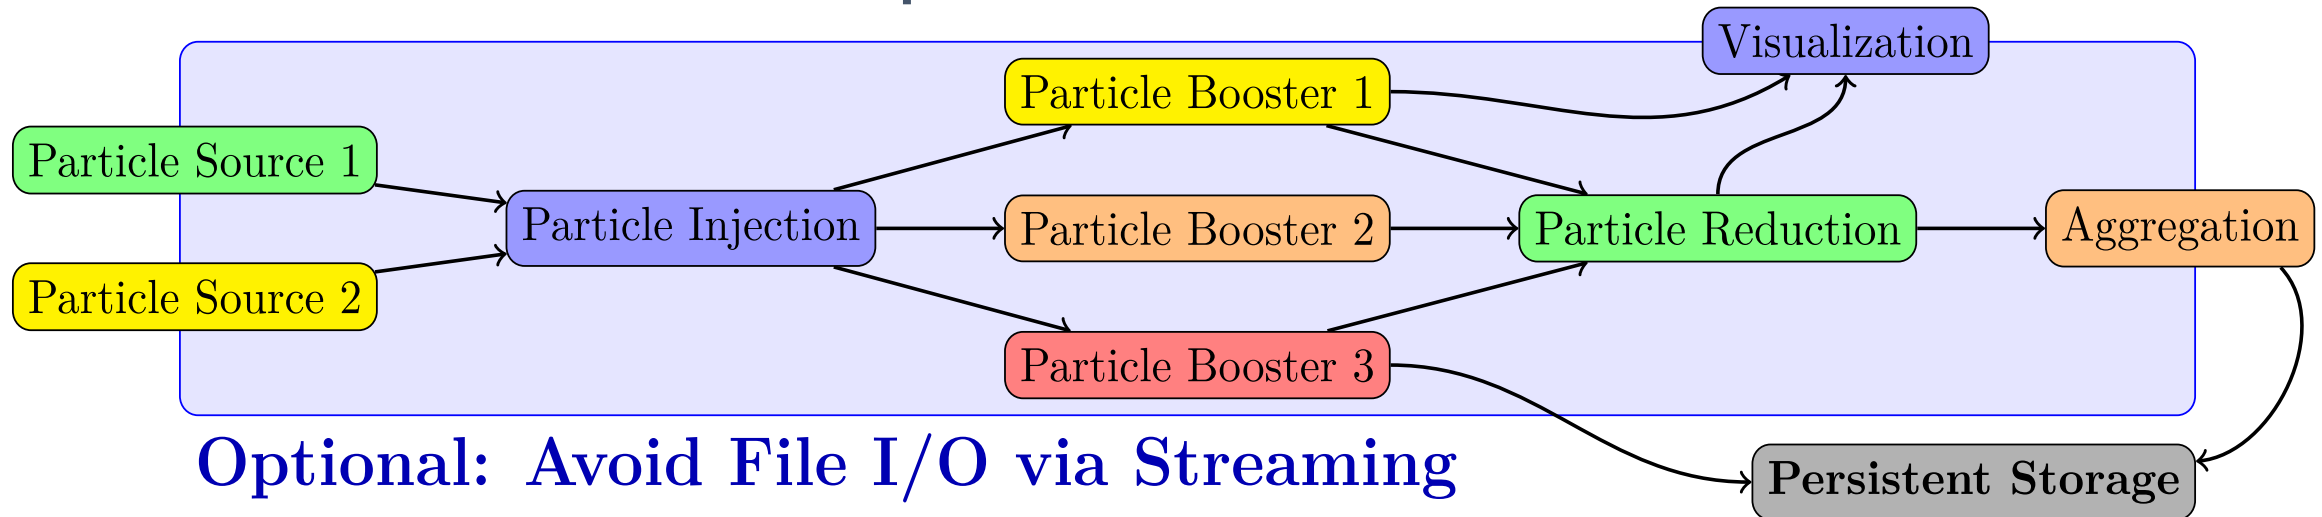


Scientific workflows are complex:




- need to span different **time** and **length scales**
- scientific modeling requires **multiple codes**, collaborating in a **data processing pipeline**
- **bridge heterogeneous models** by standardization of data


Scientific workflows are complex:




→ openPMD standard
for **particle-mesh** data
as communication layer



What is particle-mesh data?



[0:3] particles [3:6] particles [6:10] particles



Mesh

n-dimensional space,
divided into discrete cells

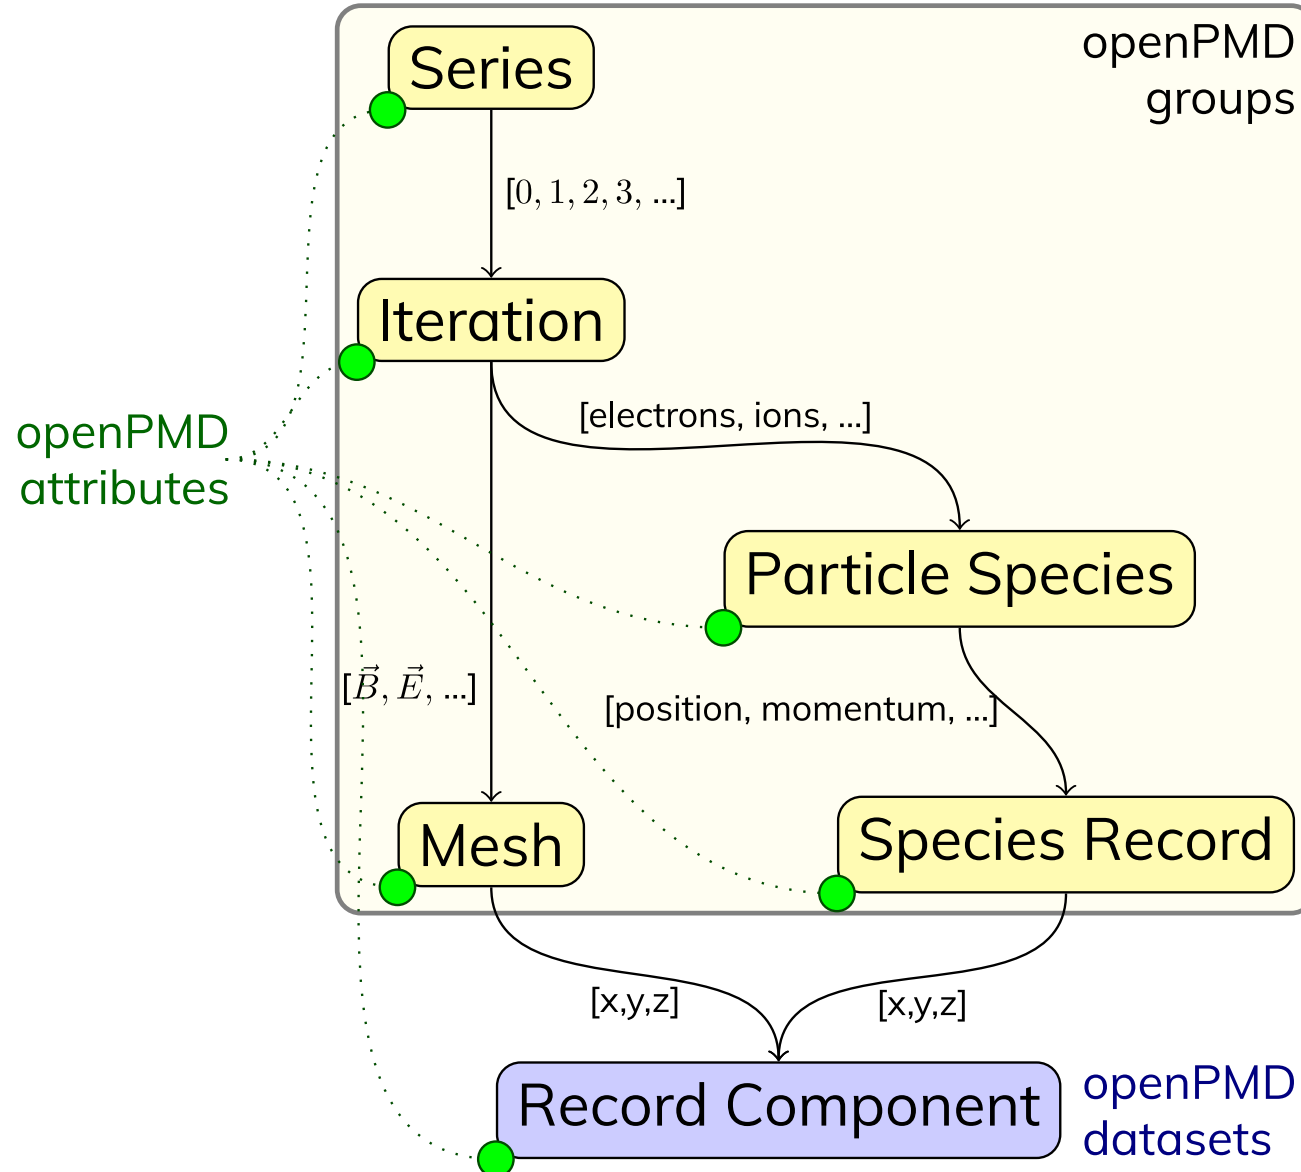
- e.g. temperature:
store a scalar number per cell
- e.g. electrical fields:
store a 3D vector per cell

Particles

A list of discrete objects,
located on the mesh

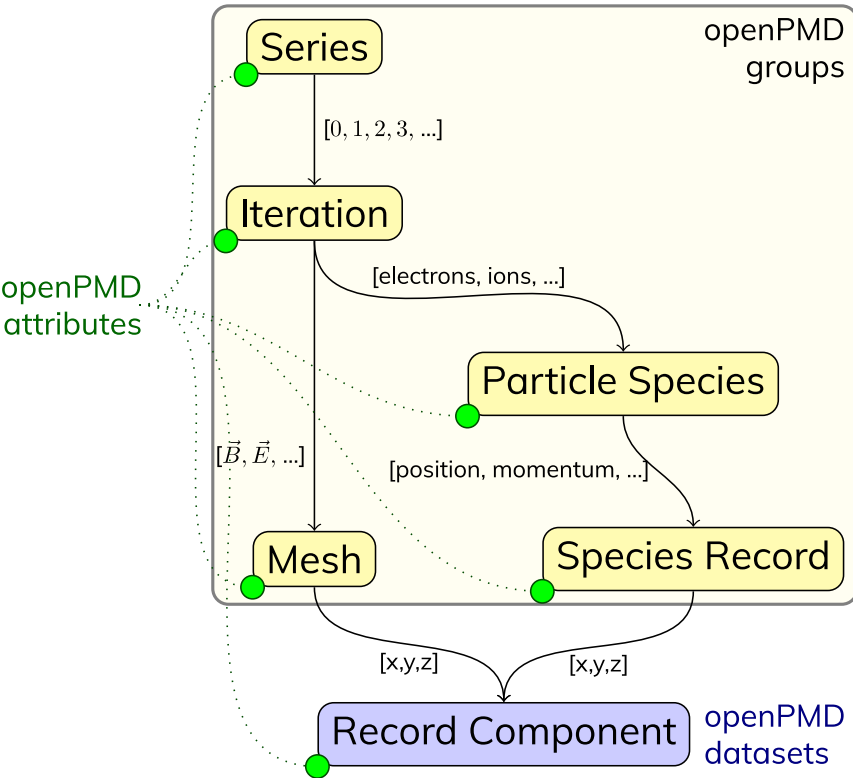
- for each particle: list its position
- optionally: list charge, weight, ...

openPMD hierarchy

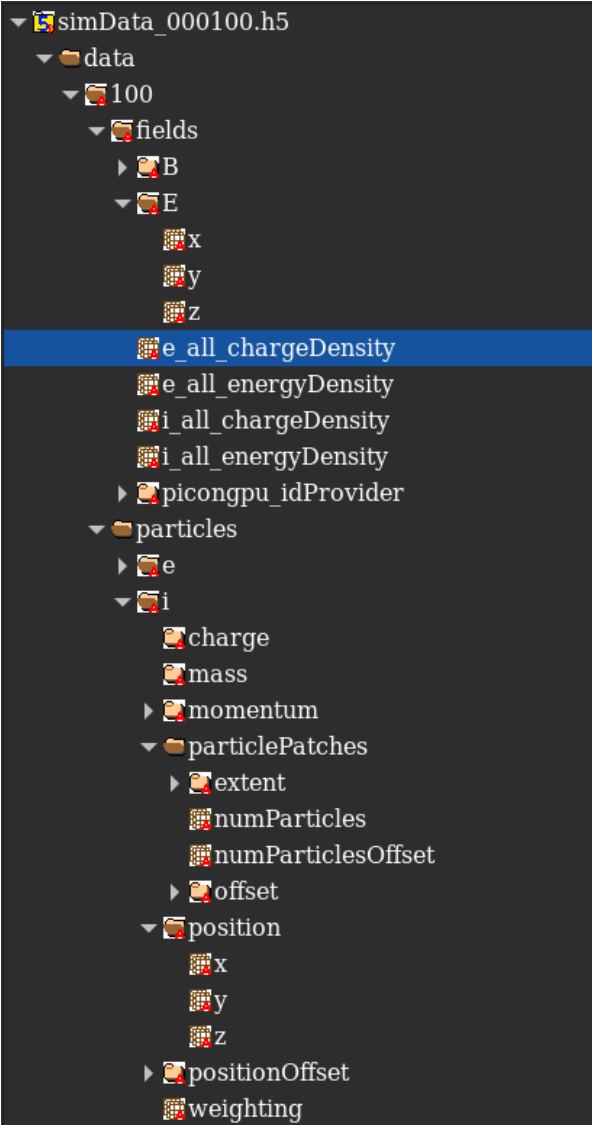


- **Structure** for series & snapshots encoded as either:
 - **files** (one file per iteration)
 - **groups** (reuse files)
 - **variables** (reuse files & variables in ADIOS2)
- Records for **physical observables** constants, mixed precision, complex numbers
- **Attributes**: unit conversion, description, relations, mesh geometry, authors, env. info, ...

Example dataset: HDF5 backend



Sample data
created with PIconGPU



File structure (HDF5):

- simData_000100.h5
 - data
 - 100
 - fields
 - B
 - E
 - x
 - y
 - z
 - e_all_chargeDensity**
 - e_all_energyDensity
 - i_all_chargeDensity
 - i_all_energyDensity

Object Attribute Info General Object Info

Attribute Creation Order: Creation Order NOT Tracked

Number of attributes = 11

| Name | Type | Array Size | Value[50](...) |
|------------------|--------|------------|------------------------------------|
| axisLabels | String | 3 | z, y, x |
| dataOrder | String | Scalar | C |
| fieldSmoothing | String | Scalar | none |
| geometry | String | Scalar | cartesian |
| gridGlobalOffset | 64-bit | 3 | 0.0, 0.0, 0.0 |
| gridSpacing | 32-bit | 3 | 1.7416798, 1.7416798, 1.7416798 |
| gridUnitSI | 64-bit | Scalar | 5.3662849982E-8 |
| position | 32-bit | 3 | 0.0, 0.0, 0.0 |
| timeOffset | 32-bit | Scalar | 0.0 |
| unitDimension | 64-bit | 7 | -3.0, 0.0, 1.0, 1.0, 0.0, 0.0, 0.0 |
| unitSI | 64-bit | Scalar | 338590.78364382515 |

Example dataset: ADIOS2 backend

| | | |
|---------|---------------------------------------|-----------------|
| float | /data/50/fields/E/x | {128, 128, 128} |
| float | /data/50/fields/E/y | {128, 128, 128} |
| float | /data/50/fields/E/z | {128, 128, 128} |
| float | /data/50/particles/e/position/x | {50053105} |
| float | /data/50/particles/e/position/y | {50053105} |
| float | /data/50/particles/e/position/z | {50053105} |
| int32_t | /data/50/particles/e/positionOffset/x | {50053105} |
| int32_t | /data/50/particles/e/positionOffset/y | {50053105} |
| int32_t | /data/50/particles/e/positionOffset/z | {50053105} |

**Hierarchical
data organization**

***n*-dim. datasets
for heavyweight data**

| | | | |
|--------|------------------------------------|------|-------------------------------|
| string | /data/50/fields/E/axisLabels | attr | = {"z", "y", "x"} |
| string | /data/50/fields/E/dataOrder | attr | = "C" |
| string | /data/50/fields/E/fieldSmoothing | attr | = "none" |
| string | /data/50/fields/E/geometry | attr | = "cartesian" |
| double | /data/50/fields/E/gridGlobalOffset | attr | = {0, 0, 0} |
| float | /data/50/fields/E/gridSpacing | attr | = {1.74168, 1.74168, 1.74168} |
| double | /data/50/fields/E/gridUnitSI | attr | = 5.36628e-08 |
| float | /data/50/fields/E/timeOffset | attr | = 0 |
| double | /data/50/fields/E/unitDimension | attr | = {1, 1, -3, -1, 0, 0, 0} |
| float | /data/50/fields/E/x/position | attr | = {0.5, 0, 0} |
| double | /data/50/fields/E/x/unitSI | attr | = 9.5224e+12 |
| float | /data/50/fields/E/y/position | attr | = {0, 0.5, 0} |
| double | /data/50/fields/E/y/unitSI | attr | = 9.5224e+12 |
| float | /data/50/fields/E/z/position | attr | = {0, 0, 0.5} |
| double | /data/50/fields/E/z/unitSI | attr | = 9.5224e+12 |

**Attributes
for self-description**

Example dataset: JSON/TOML backend

```
{
  "attributes": {
    "author": {
      "datatype": "STRING",
      "value": "franz"
    },
    "date": {
      "datatype": "STRING",
      "value": "2020-10-08 19:29:13 +0200"
    },
    "some more...": null
  },
  "data": {
    "0": {
      "attributes": {
        "cell_depth": {
          "datatype": "DOUBLE",
          "value": 4.252342224121094
        },
        "cell_height": {
          "datatype": "DOUBLE",
          "value": 1.0630855560302734
        },
        "cell_width": {
          "datatype": "DOUBLE",
          "value": 4.252342224121094
        },
        "many many more": null
      },
      "fields": {
        "B": {
          "attributes": {
            "axisLabels": {
              "datatype": "VEC_STRING",
              "value": [
                "z",
                "y",
                "x"
              ]
            },
            "position": {
              "datatype": "VEC_DOUBLE",
              "value": [
                0,
                0.5,
                0.5
              ]
            },
            "unitSI": {
              "datatype": "DOUBLE",
              "value": 40903.82224060171
            }
          },
          "data": [
            [
              "multidimensional dataset here"
            ]
          ]
        }
      }
    }
  }
}
```

- Part of the package: No need to install 3rd-party dependencies
- Useful for debugging and prototyping
- Limited parallel support
- Courtesy to Nils Lohmann's JSON library for C++
- With recent release: Convert output to TOML
Idea: openPMD formatted configuration files

Reference Implementation in C++ & Bindings: Python and Julia

Online Documentation:
openpmd-api.readthedocs.io

Open-Source Development & Tests:
github.com/openPMD/openPMD-api

INSTALLATION

Installation

Changelog

Upgrade Guide

USAGE

Concepts

First Write

Include / Import

Open

Iteration

Attributes

Data

Record

Units

Register Chunk

Flush Chunk

After successful installation, you can start using openPMD-api as follows:

C++17

```
#include <openPMD/openPMD.hpp>

// example: data handling
#include <numeric> // std::iota
#include <vector>   // std::vector

namespace io = openPMD;
```

Python

```
import openpmd_api as io

# example: data handling
import numpy as np
```

Open


Write into a new openPMD series in `myOutput/data_<00...N>.h5`. Further file formats than `.h5` (HDF5) are supported: `.bp` (ADIOS1/ADIOS2) or `.json` (JSON).

```
auto series = io::Series(
    "myOutput/data_005T.h5",
    io::Access::CREATE);
```

```
series = io.Series(
    "myOutput/data_005T.h5",
    io.Access.create)
```

| | | |
|---|---|---|
| ✓ | All checks have passed | 25 successful checks |
| ✓ | macOS / appleclang12_py_mpi_h5_ad2 (pull_request) | Successful in 17m Details |
| ✓ | Windows / MSVC w/o MPI (pull_request) | Successful in 6m Details |
| ✓ | Intel / ICC C++ only (pull_request) | Successful in 7m Details |
| ✓ | Tooling / Clang ASAN UBSAN (pull_request) | Successful in 58m Details |
| ✓ | Nvidia / CTK@11.2 (pull_request) | Successful in 4m Details |
| ✓ | Linux / clang8 py38 mpich h5 ad1 ad2 newLayout (pull request) | Successful in 29m Details |

Rapid and easy installation on any platform:



```
python3 -m pip install
openpmd-api
```



```
brew tap openpmd/openpmd
brew install openpmd-api
```



```
cmake -S . -B build
cmake --build build
--target install
```



```
conda install
-c conda-forge
openpmd-api
```



```
spack install
openpmd-api
```



```
module load openpmd-api
```

A Huebl, F Poeschel, F Koller, J Gu, et al.
"openPMD-api: C++ & Python API for Scientific I/O with openPMD" (2018) [DOI:10.14278/rodare.27](https://doi.org/10.14278/rodare.27)

Hands-On:

openPMD-api: basic object model

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Read the TODO comments inside `src/openPMDOutput.hpp`:

unitDimension


automated description of physical dimension
only powers of base dimensions

length **L**, mass **M**, time **T**, electric current **I**,
thermodynamic temperature **theta**,
amount of substance **N**, luminous intensity **J**

Magnetic field: $[B] = M / (I * T^2)$
→ (0, 1, -2, -1, 0, 0, 0)

unitSI (recommended)

relation to an absolute unit system




Wikimedia Commons

Findable: Standardized metadata to identify the data producer

```
string    /author      attr    = "franz"  
string    /software    attr    = "PIConGPU"  
string    /softwareVersion attr    = "0.5.0-dev"
```

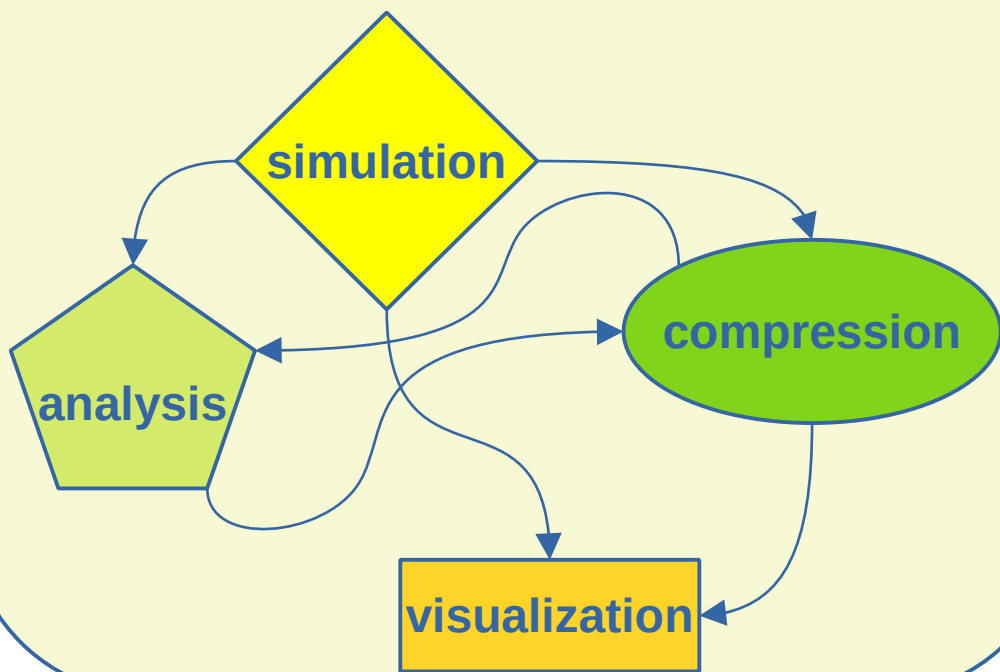
Accessible: Open standard, implementable in various formats



*currently implemented,
but not limited to

Interoperable:

Data exchange spans applications, platforms and teams



Reusable:


Rich and standardized description for physical quantities

| Name | Value |
|------------------|--------------------------------|
| axisLabels | [b'z' b'y' b'x'] |
| dataOrder | b'C' |
| fieldSmoothing | b'none' |
| geometry | b'cartesian' |
| gridGlobalOffset | [0. 0. 0.] |
| gridSpacing | [4.252342 1.0630856 4.252342] |
| gridUnitSI | 4.1671151662e-08 |
| position | [0. 0. 0.] |
| timeOffset | 0.0 |
| unitDimension | [-3. 0. 1. 1. 0. 0. 0.] |
| unitSI | 15399437.98944343 |

“The FAIR Guiding Principles for scientific data management and stewardship” (Mark D. Wilkinson et al.)

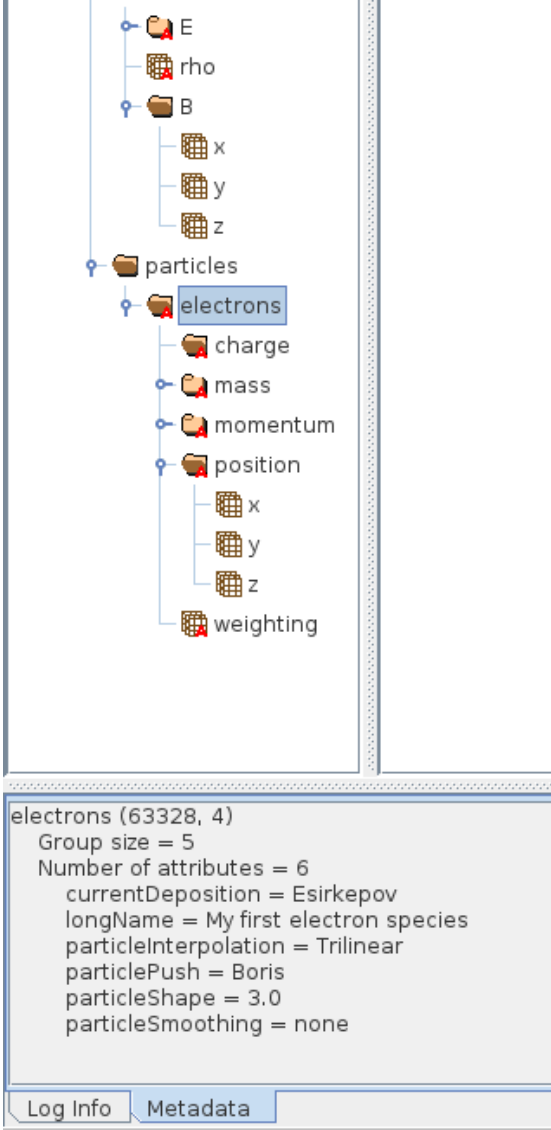
Extensions: e.g. ED-PIC

Field image → field solver, smoothing



similar:

Emittance → particle push, field solver, shape



Hands-On:

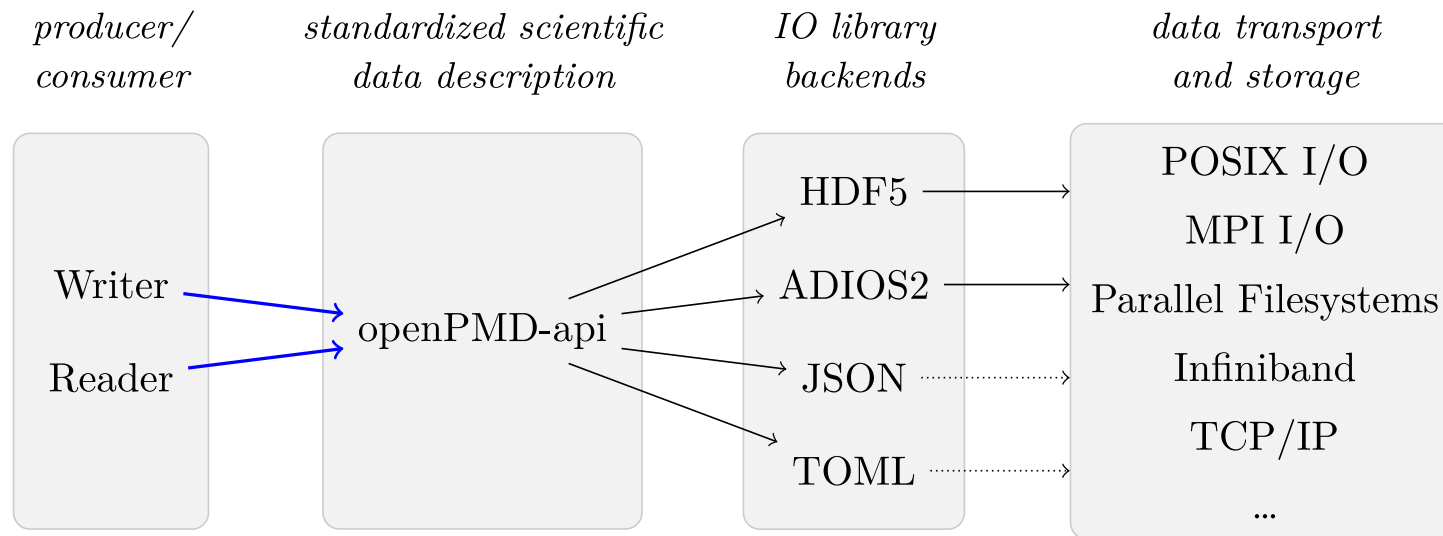
openPMD-api: metadata

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Read the TODO comments inside `src/openPMDOutput.hpp`:

openPMD-api – open stack for scientific I/O



- MPI support at all levels
- Implemented in C++17
- Bindings in C++17, Python and (dev version only) Julia
- Specify backend at runtime:
I/O library, transport, compression,
streaming, aggregation, ...

```
import openpmd_api as io

# pick and configure backend via JSON/TOML or inferred from filename extension
adios_config = """
    backend = "adios2"
    [[adios2.dataset.operators]]
    type = "blosc" # activate compression
"""

mode = io.Access.create
series = io.Series("simOutput.h5", mode,
    """{"hdf5": {"vfd": {"type": "subfilings"}}}""")
series = io.Series("simOutput.bp5", mode, adios_config)
series = io.Series("simOutput.sst", mode, "@./or/load/config/from/file.json")
series = io.Series("simOutput.json", mode)
```


Hands-On:

openPMD-api: visualization, backend configuration


Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the instructions inside `src/next_steps.md`:

I/O Performance lags behind Compute Performance



| | Titan | | Summit | | Frontier | |
|-------------------|-------|----------|--------|----------|----------|----------|
| Peak Performance: | 27 | Pflop/s | 200 | Pflop/s | 1.6 | Eflop/s |
| FS Throughput: | 1 | TiByte/s | 2.5 | TiByte/s | 5~10 | TiByte/s |
| FS Capacity: | 27 | PiByte | 250 | PiByte | 500~1000 | PiByte |

| Growth Factor | |
|---------------|-------|
| | ~60 |
| | 5~10 |
| | 18~37 |

- **parallel bandwidth** insufficient for HPC at full scale
- **filesystem capacity** insufficient for HPC at full scale

Same trend in **experiments?**

- Increasing **camera resolutions and data rates**

Compute Performance Outpaces Storage Performance



Titan

Peak Performance: 27 Pflop/s
FS Throughput: 1 TiByte/s
FS Capacity: 27 PiByte



Summit

200 Pflop/s
2.5 TiByte/s
250 PiByte



Frontier

1.6 Eflop/s
5~10 TiByte/s
500~1000 PiByte

Growth Factor

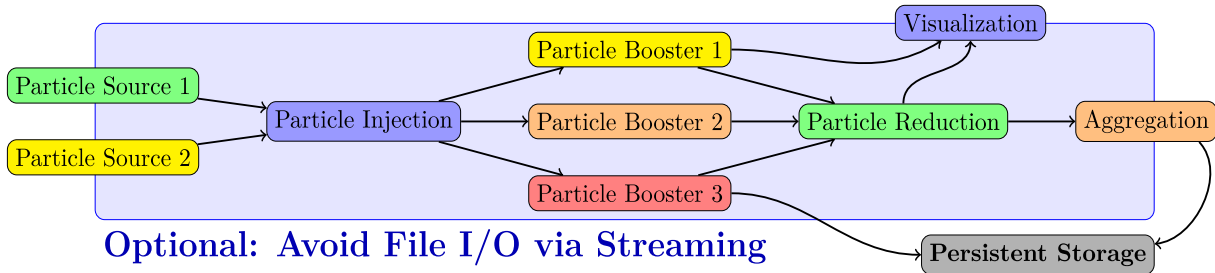
~60

5~10


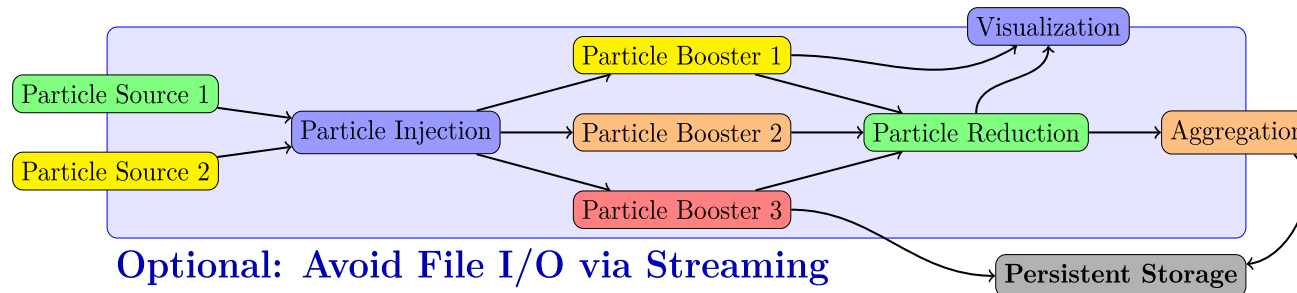
18~37

Why does this concern us?

- Heterogeneous data processing pipelines traditionally have large I/O usage
- Scalable alternative: Streaming




Streaming: Don't touch the Filesystem at all



Challenge:
Compute a **balanced, aligned, local** mapping between two applications that remains useful in the problem domain

Break through Filesystem Bandwidth with Streaming

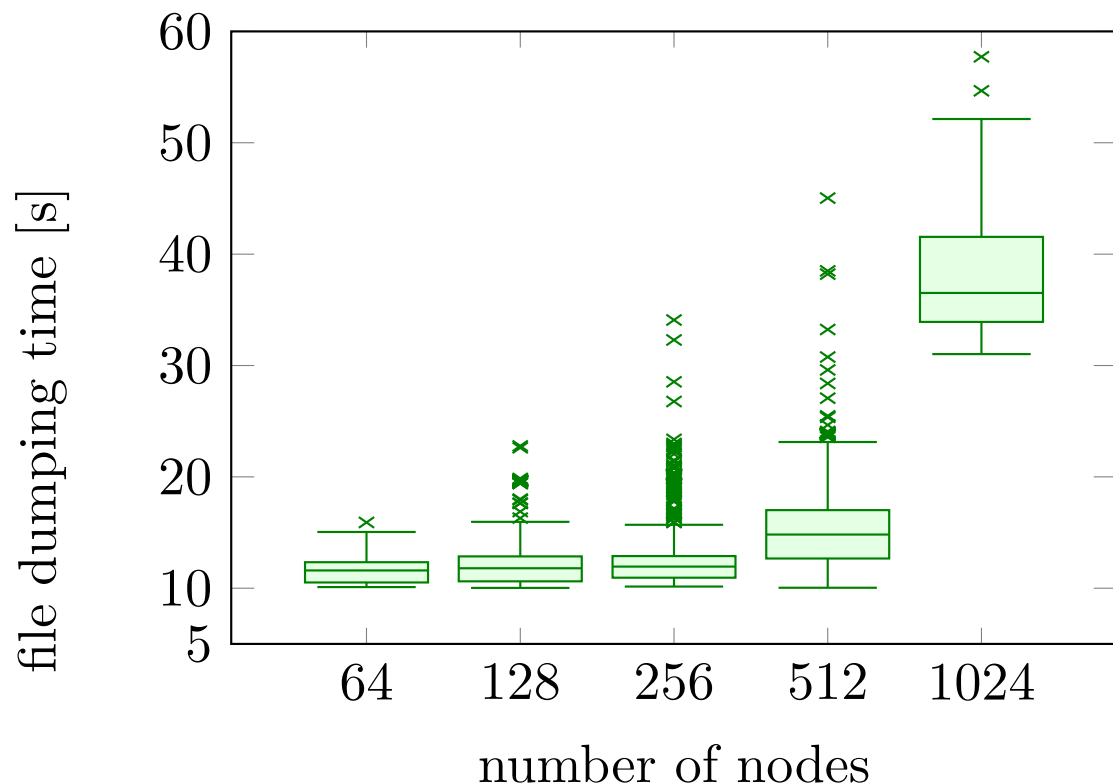


Memory-bound simulations reach the I/O system limits at a fraction of full scale

- Summit FS bandwidth (2.5TiByte/s) reached at 512 nodes (~11% of system size)
- Streaming workflows unaffected by filesystem bandwidth, use Infiniband hardware to scale beyond it

(benchmarks at 1024 nodes done after Summit system upgrade)

Summit: Performance fluctuations on single ranks




Same results, different display:

- Plot every single measurement
- Visualize reproducibility
- Box: 50% of measurement points
- Whiskers: “normal” measurements
- Others: outliers

Evaluation for file-only setup:

- Median time slightly raised at 512 nodes
Scaling stops due to PFS limit after that
- Outliers increasing with scale
- Outliers fatal in parallel contexts




Same results, different display:

- Plot every single measurement
- Visualize reproducibility
- Box: 50% of measurement points
- Whiskers: “normal” measurements
- Others: outliers

stream+file setup (stream part):

- Overall times are lower
- Median between 5 and 7 seconds
- Outliers less dominating by far

For good throughput: Local streaming patterns, Infiniband/RDMA



Local streaming:

Distribute data chunks only within a node

Non-local streaming:


Distribute data chunks globally, optimize for balance and alignment

Straight lines:

Infiniband/RDMA

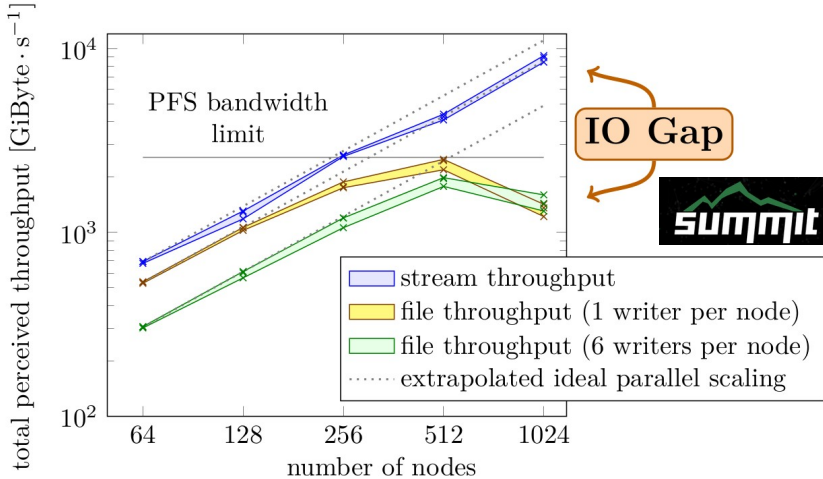
Dashed lines:

TCP/sockets



Setup: Couple PIconGPU with a scattering code (GAPD) exchange particle data only


Performance: Data Layouts and no-file I/O



Streaming Data Pipelines:

[DOI:10.1007/978-3-030-96498-6_6](https://doi.org/10.1007/978-3-030-96498-6_6)

by F Poeschel, A Huebl et al., SMC21 (2022)




Fast Compressors Needed:

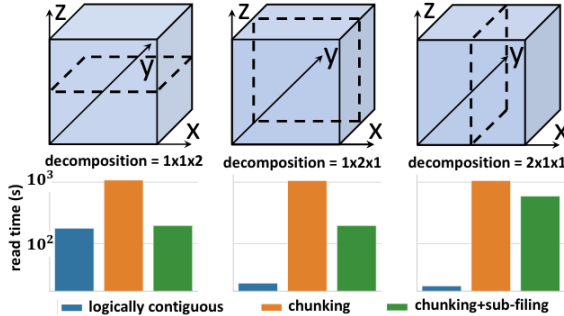

[DOI:10.1007/978-3-319-67630-2_2](https://doi.org/10.1007/978-3-319-67630-2_2)

by A Huebl et al., ISC DRBSD-1 (2017)

ADIOS openPMD-api w/ WarpX



>5.5TB/s FS BW: two-tier lustre w/ high-performance storage & progressive files



Impact of decomposition schemes when reading

Online Data Layout Reorganization:

[DOI:10.1109/TPDS.2021.3100784](https://doi.org/10.1109/TPDS.2021.3100784)

by L Wan, A Huebl et al., TPDS (2021)

Hands-On:

openPMD-api: streaming I/O

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the instructions inside `src/visualize.py` (open as a Jupyter Notebook):

openPMD powered Projects and Users



Documents:

- **openPMD standard** (1.0.0, 1.0.1, 1.1.0)
the underlying file markup and definition
A Huebl et al., doi: 10.5281/zenodo.33624

Language Binding:

- **openPMD-api** (HZDR, CASUS, LBNL)
reference API for openPMD data handling
maintainers: A Huebl, J Gu, F Poeschel et al.

HiPACE++ → VisualPIC
Credit: M.Thévenet
& A. Ferran Pousa (DESY)



MALA → ParaView
Credit: A. Cangi (CASUS)

- **Wake-T** (DESY)
fast particle-tracking code for plasma-based accelerators
maintainer: A Ferran Pousa
- **HiPACE++** (DESY, LBNL)
3D GPU-capable quasi-static PIC code for plasma accel.
maintainers: M Thevenet, S Diederichs, A Huebl
- **Bmad** (Cornell)
library for charged-particle dynamics simulations
maintainers: D Sagan et al.
- **MALA** (CASUS, SNL)
ML models that replace DFT calculations in materials science
maintainers: Attila Cangi & Sivasankaran Rajamanickam
- and more...

see also: <https://github.com/openPMD/openPMD-projects>

Analysis and Visualization

```
In [1]: import numpy as np
import matplotlib notebook
# or '%matplotlib inline' for non-interactive plots
# or '%matplotlib widget' when using JupyterLab (github.com/matplotlib/jupyter-matplotlib)
import matplotlib.pyplot as plt
from openpmd_viewer import OpenPMDTimeSeries
```

```
In [2]: # Replace the string below, to point to your data
ts = OpenPMDTimeSeries('/home/franzpoeschel/singularity_build/pic_run/openPMD')
```

```
In [3]: # Interactive GUI
ts.slider()
```

-

+

iteration

1900

Field type

Field:

B

E

e_all_chargeDensity

e_all_energyDensity

picongpu_idProvider

Coord:

x

y

z


Particle quantities

Particle selection

Plotting options

Always refresh

Refresh now!



openPMD/openPMD-viewer



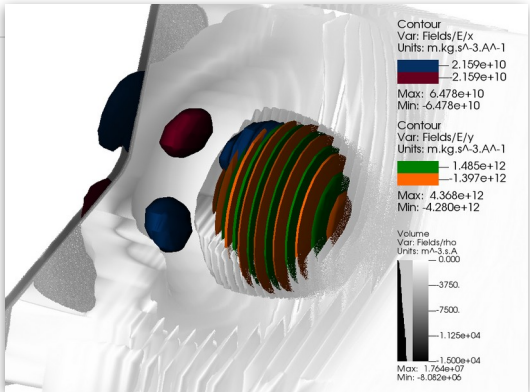
Standardization of data
→ integration into modern scientific compute workflows



DASK



ParaView



Hands-On:

openPMD-viewer: visualizing data written by a PIConGPU LaserWakefield simulation

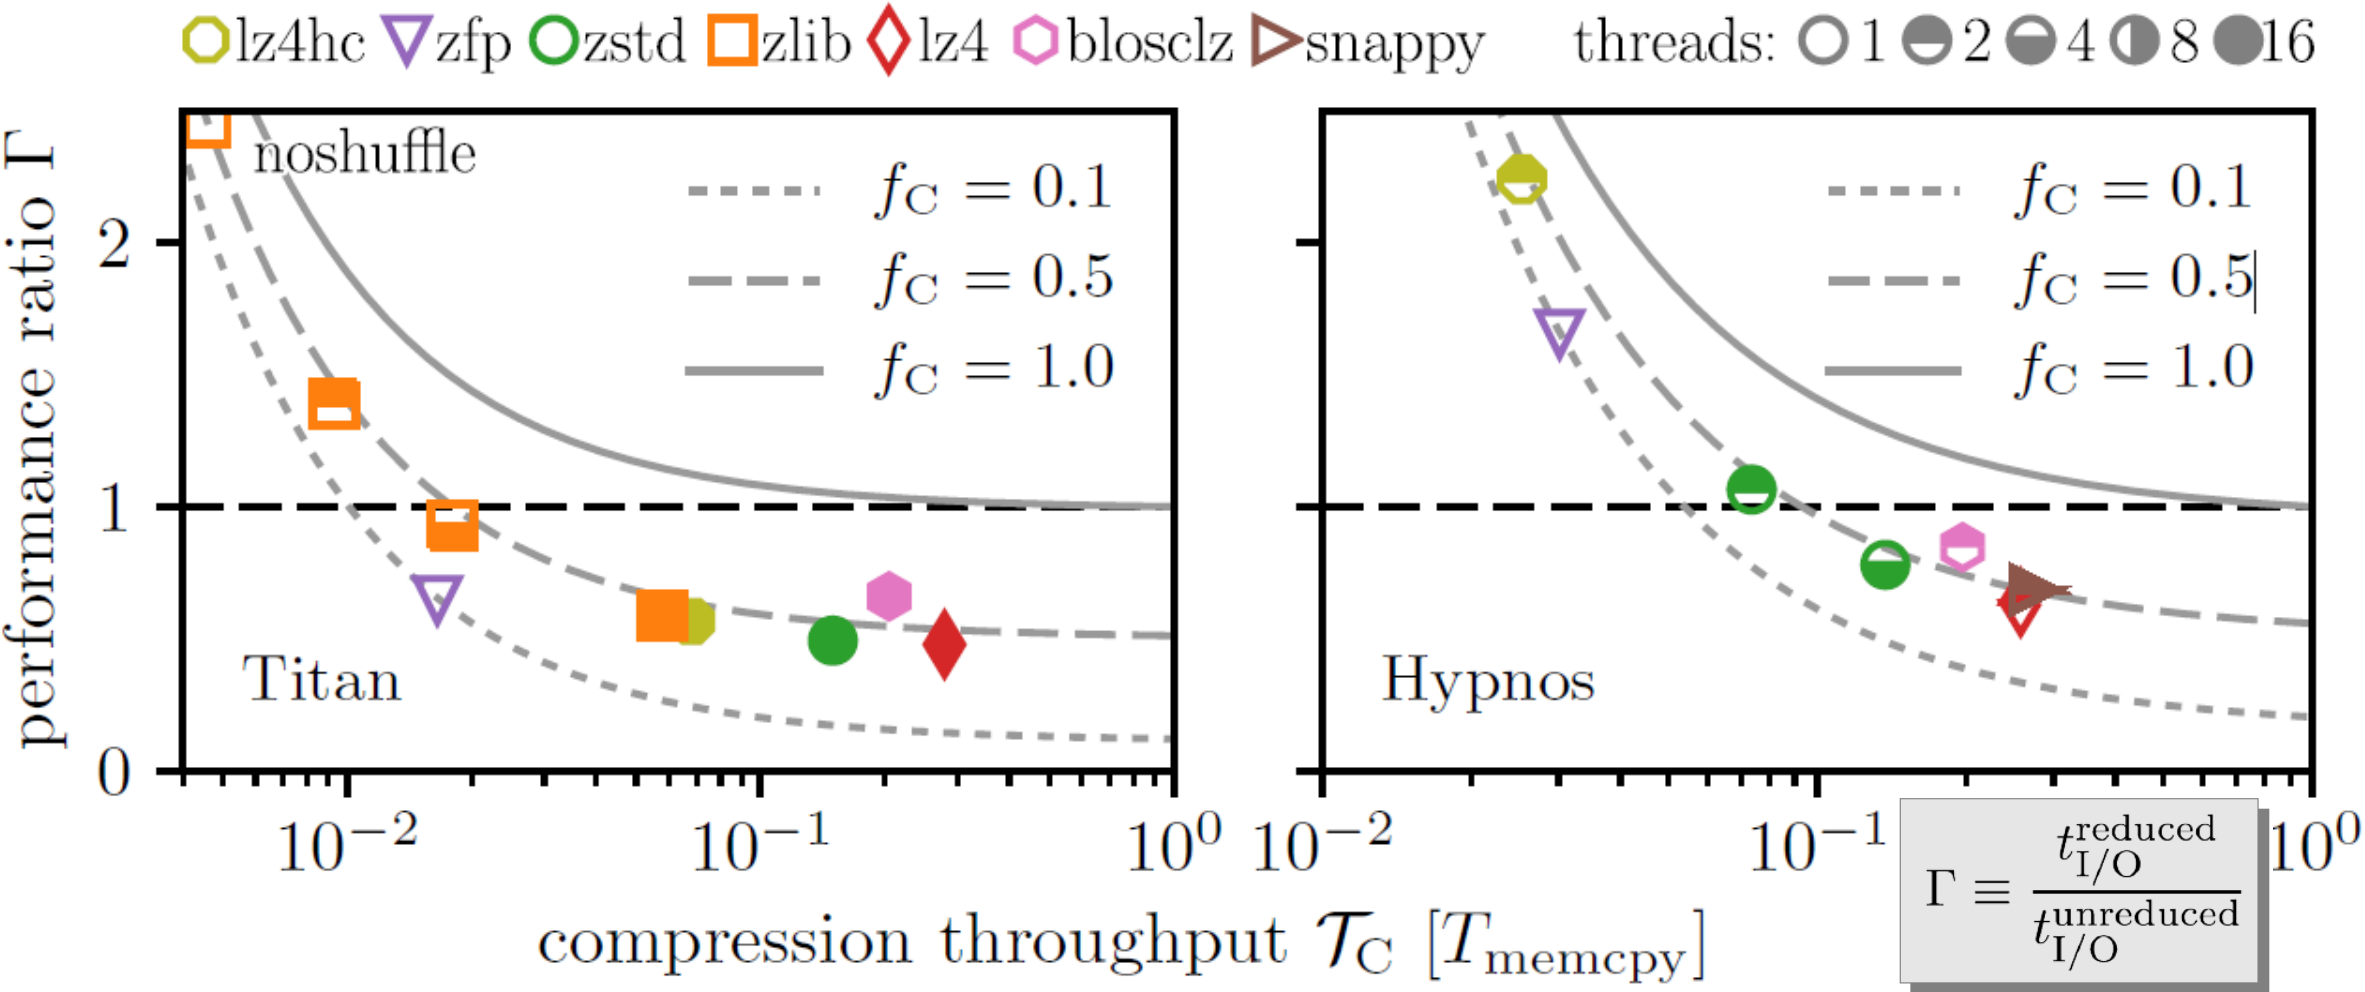
Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the instructions inside `next_steps.md`:

Just compress the data and everything will be fine ...?



A. Huebl et al., "On the Scalability of Data Reduction Techniques in Current and Upcoming HPC Systems from an Application Perspective", In: Lect. Notes Comput. Sci. 10524.4, pp.15-20 (2017)

Hands-On: **compression**

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at `https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop`

Launch Notebook at `https://www.lumi.csc.fi/`

Continue with the instructions in the openPMD-viewer Notebook

TBD: Data workflow in a realistic application

Advanced Hands-On:

openPMD-api: Span API

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at `https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop`

Launch Notebook at `https://www.lumi.csc.fi/`

Read the instructions inside `src/openPMDOutput.hpp`

TBD: Recap constant components

Advanced Hands-On:

openPMD-api:

Python API: constant components

Module environment at `/project/project_465001310/workshop_software/env.sh`

Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the TODO comments inside `write_parallel.py`

TBD: Parallel I/O, parallel file systems

Advanced Hands-On:

openPMD-api:

Python API: Parallel Writing

Module environment at `/project/project_465001310/workshop_software/env.sh`


Materials at https://github.com/alpaka-group/alpaka-workshop-slides/tree/oct2024_workshop

Launch Notebook at <https://www.lumi.csc.fi/>

Read the TODO comments inside `write_parallel.py`

Where to find us

→ head to <https://github.com/openPMD/>



openPMD
Open Standard for Particle-Mesh Data Files
<https://www.openPMD.org> axelhuebl@lbl.gov

Repositories 17


Packages


People 50


Teams 5


Projects


Pinned repositories


**openPMD-standard**
Open Standard for Particle-Mesh Data Files
☆ 41 🔗 17

**openPMD-projects**
Overview on Projects around openPMD
☆ 4 🔗 4


**openPMD-viewer**
Python visualization tools for openPMD files
Jupyter Notebook ☆ 35 🔗 26

**openPMD-api**
C++ & Python API for Scientific I/O
C++ ☆ 55 🔗 30

**openPMD-visit-plugin**
Plugin allowing VisIt to read openPMD files
C ☆ 8 🔗 3

**openPMD-example-datasets**
HDF5 Example Files
Python ☆ 5 🔗 1

...and of course <https://openpmd-api.readthedocs.io/>



The **openPMD standard** is co-authored by [Axel Huebl](#), [Rémi Lehe](#), Jean-Luc Vay, David P. Grote, Ivo F. Sbalzarini, Stephan Kuschel, David Sagan, Frédéric Pérez, Fabian Koller, [Franz Poeschel](#), Carsten Fortmann-Grote, Ángel Ferran Pousa, Juncheng E, [Maxence Thévenet](#), and Michael Bussmann.

The authors are thankful for the **community contributions** to libraries, software ecosystem, user support, review and integrations. Particularly, thank you to Yaser Afshar, Lígia Diana Amorim, James Amundson, Weiming An, Igor Andriyash, Ksenia Bastrakova, [Jean Luca Bez](#), Richard Briggs, Heiko Bureau, Jong Choi, Ray Donnelly, Dmitry Ganyushin, Marco Garten, Lixin Ge, Berk Geveci, Daniel Grassinger, Alexander Grund, [Junmin Gu](#), Marc W. Guetg, Ulrik Günther, Sören Jalas, Manuel Kirchen, John Kirkham, Scott Klasky, Noah Klemm, Fabian Koller, Mathieu Lobet, Christopher Mayes, Ritiek Malhotra, Paweł Ordyna, Richard Pausch, [Norbert Podhorszki](#), David Pugmire, Felix Schmitt, [Erik Schnetter](#), Dominik Stańczak, Klaus Steiniger, Michael Sippel, Frank Tsung, Lipeng Wan, René Widera, and Erik Zenker!